

KEYWORDS: Chatbot, Retrieval-Augmented Generation, Hybrid Retrieval, Academic Advising, Hallucination Detection, Code-Mixed NLP, Natural Language Processing, Educational Technology, Information Retrieval.

Abstract: We extend a BM25 + dense hybrid retrieval-augmented generation (RAG) chatbot for university academic advising into an adaptive pipeline with four independently-ablated components: (i) domain-fine-tuned bilingual sentence embeddings, trained via contrastive (question, answer) pairs on a strict train/validation/test split; (ii) per-query-type fusion routing, which sends entity-heavy queries (e.g., course-code and faculty lookups) through a heavily lexical-weighted linear blend and open-ended queries through the original, more balanced linear BM25/vector blend – entity-heavy queries were originally routed through Reciprocal Rank Fusion instead, until a deconfounding diagnostic found that specific choice measurably negative with no offsetting benefit and we replaced it, verified live to change no real-world behavior while narrowing a residual ranking-quality gap; (iii) prerequisite-graph augmentation, which answers multi-hop prerequisite-chain questions by traversing a directed graph built from the knowledge base’s own edge data, a query shape neither lexical nor dense retrieval can answer from a single chunk; and (iv) confidence-gated abstention, calibrated per route against a labeled out-of-scope subset (augmented with database-verified synthetic examples), with a sufficient-context override that prevents the system from declining to answer when direct structural evidence is already in hand. We evaluate on 200 English and 100 Banglish (Bengali-English code-mixed) test queries with paired significance testing and standard IR metrics (Recall@k, MRR, nDCG@10), and report every result obtained, including negative ones: a cross-encoder reranker – both generic and, after we caught and corrected a training/deployment mismatch, properly fine-tuned on the corpus’s own hard negatives – was found to significantly *hurt* the pipeline in both cases and is disabled by default; BM25-level Banglish spelling-variant normalization shows a null effect on the main Banglish test set. Cross-lingual query translation also showed a null effect on that same test set, but we traced this to a structural blind spot in the test set itself and found translation still helps (61.5% to 84.6% top-5 accuracy, $n = 13$, a real replication of the same direction at a still-small but growing sample) on a dedicated stress test built for the scenario it targets, though a larger, noisier expanded version of that same stress test no longer reaches even suggestive statistical support after this revision’s other retriever improvements partly closed the gap it was designed to detect. Under the current retriever (hard-negative-mined embeddings, an exact-match coverage fix, and a BM25 k_1/b retuning), the full adaptive pipeline is in full statistical parity with every single-method baseline on all four automated similarity metrics: tied with full hybrid and vector-only on all four, and not confirmed behind BM25-only on any of them either – the closest case, METEOR, is marginally significant under a paired t -test alone ($p = 0.041$) but not under a matched Wilcoxon test ($p = 0.306$), so we treat it as inconclusive rather than confirmed, consistent with how we treat other test-disagreeing results throughout this paper. This supersedes an earlier revision’s confirmed METEOR loss, which we found had gone stale after the exact-match and BM25 fixes and re-measured rather than leave standing. Parity is not automatically an advantage, and we tested rather than assumed it: a deconfounding diagnostic isolating adaptive routing’s own fusion-method choice from a separate exact-match ranking guarantee found the original choice (Reciprocal Rank Fusion for entity-heavy queries) was significantly negative on entity-heavy retrieval once that guarantee was removed, not merely neutral – the parity with BM25-only was being borrowed from that separate mechanism, not evidence the routing choice itself was beneficial. We acted on this rather than only disclose it: replaced Reciprocal Rank Fusion with linear fusion at the same weighting for entity-heavy queries, verified live that this changes no observed behavior on the current test set (the exact-match mechanism already determines every real outcome), and measured a genuine improvement in the residual ranking-quality gap as a direct, confirmed result. The system’s demonstrated value is therefore in capabilities the single-method baselines cannot perform at all, not in beating them on automated similarity metrics: abstaining on only 0.5–1.0% of queries it cannot ground (down from 10.5% before a coverage fix to the exact-match mechanism), correctly answering multi-hop prerequisite-chain questions via graph traversal, and operating over a bilingual corpus with comparable results in both languages. An automated LLM-judge faithfulness check (RAGAS-style) shows the deployed baselines score 0.92 (BM25-only 0.916, full hybrid 0.923) and the full pipeline scores 0.81 (descriptively, not a paired comparison at this level); a matched-subset paired test finds this gap is not statistically significant against the two strongest baselines, and a second, architecturally independent NLI-based faithfulness check shows only weak, near-chance-level agreement with the LLM judge at the row level – together, evidence too inconclusive to confirm a genuine grounding regression, but not clean enough to rule one out either, which we report as an open question rather than resolve in either direction. We report this system’s contribution honestly: the underlying retrieval mathematics (lexical/dense linear fusion, RRF) is established, not novel, and every generation-quality comparison against single-method baselines lands at parity rather than a win. What is genuinely new is at the mechanism and methodology level rather than the ranking-formula level – an ambiguous-entity disambiguation hint that builds a clarifying question directly from a corpus field the tied candidates disagree on, a deconfounding-diagnostic method for separating a routing decision’s own contribution from a masking downstream mechanism (used here to find and fix a real, measured regression), and a bilingual application of sufficient-context classification not reported elsewhere in the literature we surveyed – integrated and validated for an under-studied bilingual advising setting. We disclose every limitation and methodological pitfall we found along the way, including our own.

Adaptive Hybrid Retrieval with Confidence-Gated Abstention for Bilingual Academic Advising

Anonymous Author(s) , Anonymous Institute , anonymous@anonymous.edu

August 1, 2026

1 INTRODUCTION

Open-credit university systems place substantial demands on students to independently plan course sequences, interpret evolving policy documents, and schedule their studies, while mentor availability per student remains limited. Despite the growing availability of digital academic resources, students at open-credit institutions continue to report difficulty accessing timely, personalized guidance on course selection and policy questions, particularly during the first year (Nguyen et al., 2023). This gap is especially acute for first-year students, who are simultaneously adjusting to a new academic environment and to the self-directed nature of open-credit planning. At the institution this system targets, a substantial share of student questions are also not posed in formal English but in Banglish – Bengali-English code-mixed text typed in Latin script – a register general-purpose and even most retrieval-augmented university chatbots are not built or evaluated for.

Conversational agents are a natural candidate for addressing this gap: they can process natural language input and sustain multi-turn, human-like conversational exchanges with users (Brandtzaeg and Følstad, 2017), and can offer fast, personalized assistance at a scale that individual mentors cannot (Shawar and Atwell, 2007). However, general-purpose conversational agents built on large language models (LLMs) are prone to hallucination – the generation of fluent but factually incorrect content – which is particularly risky in an advising context where a single incorrect prerequisite or deadline can have real consequences for a student’s academic trajectory.

This paper substantially extends our own earlier system, Rahman et al. (2025), a hybrid BM25/ChromaDB academic-advising chatbot for the same institution that established a simple lexical/dense hybrid retriever outperforms either retrieval method alone specifically on entity-heavy queries (alphanumeric course-code identification), while blending in dense retrieval on top of a well-tuned lexical stream neither reliably helps nor hurts on open-ended natural-language questions – using a single fixed fusion weight ($\lambda = 0.5$) throughout. That system’s own limitations explicitly identify the absence of Bangla/Banglish support as a gap left to future work, and offer no confidence-gated abstention mechanism, no reasoning over the corpus’s prerequisite structure beyond retrieving whichever chunk states it, and no ablation of a reranking stage. This paper’s contribution is to take the fixed-fusion-weight finding as a design constraint rather than a conclusion, and to close each of those gaps directly rather than merely name them: if a single global fusion weight cannot be shown to beat lexical retrieval on any query type, the right response is to stop searching for one global weight and instead route different query types differently, add the structured signals (a prerequisite graph, exact metadata matches) that text retrieval alone cannot recover, add a mechanism for declining to answer when none of this is enough, and extend the system to the bilingual setting its predecessor’s own limitations called for – rather than accept a single fixed architecture as final. The dataset itself has also been substantially revised for this paper (Section 3.1), not merely re-scored with the predecessor’s original data. Concretely, this paper:

- domain-adapts the retriever’s sentence embeddings via contrastive fine-tuning on the corpus’s own (question, answer) pairs, under a strict split that never trains on validation or test data, and reports the held-out accuracy gain this produces before any other component is touched;
- replaces the single fixed- λ fusion rule with adaptive routing by query type (entity-heavy vs. open-ended), motivated directly by the entity/open-ended asymmetry the prior ablation already established;
- adds a directed prerequisite graph, built from the knowledge base’s own edge data, to correctly answer multi-hop prerequisite-chain questions that no single retrieved chunk contains;
- adds a confidence-gated abstention mechanism, calibrated per route against labeled out-of-scope examples, together with a sufficient-context override that prevents abstaining when direct structural evidence for an answer is already available – both signals computed as a side effect of retrieval already performed, at no additional model call;
- adds an ambiguous-entity disambiguation mechanism that, when a named entity resolves to more than one real record, builds its clarifying question around a field the candidates actually disagree on rather than asking the user to simply retype a full name, validated at $n = 220$ against both a plain notice and an earlier, less directive version of the same idea (all three judged criteria, $p < 0.0001$);
- extends the corpus and every evaluation to a second language variety, Banglish, and reports the corresponding results and ablations (query translation, code-mixed lexical normalization) with the same rigor, including where they do not help; and
- reports every ablation’s outcome, including the negative ones (a non-fine-tuned cross-encoder reranker, disabled by default) and every methodological limitation found along the way, since the standing requirement for this work is that no reported number be anything other than the output of code that was actually run.

Accordingly, we ask three research questions: (RQ1) does replacing single-method or single- λ retrieval with the adaptive pipeline described above improve automated semantic-quality scores, and if not, what does it add instead;

(RQ2) do the individual novel components (fine-tuned embeddings, adaptive routing, graph augmentation, abstention, reranking) each help, hurt, or have no measurable effect, evaluated by direct ablation; and (RQ3) does the system perform comparably on Banglish queries, and do techniques specifically aimed at the code-mixed setting (query translation, lexical normalization) measurably help on this corpus. Sections 2–4 address these questions in turn, and Section 5 discusses their implications together with the system’s limitations.

2 RELATED WORK

2.1 Chatbots in Education

Conversational agents have been applied to education since ELIZA, an early system based on pattern matching and scripted responses (Adamopoulou and Moussiades, 2021). Subsequent work has examined chatbots for FAQ-style university support (El-Ashmawi et al., 2023), knowledge-base-driven question answering (Formica et al., 2023), and retention-oriented advising support, such as the Lilo system, which combines advising with social-emotional support to address “summer melt” among newly admitted students (Nguyen et al., 2023). Hwang and Chang (2021) survey the opportunities and open challenges of educational chatbots, noting that evidence of effectiveness across diverse educational settings remains limited. More recently, Swacha and Gracel (2025) survey 47 RAG-based educational chatbots and identify hallucination mitigation as the primary motivation for adopting RAG in this domain, while noting that direct hallucination measurement, rather than similarity-based proxies, remains rare across the surveyed literature. None of the systems surveyed report evaluation on code-mixed or bilingual student input, which is the gap Section 4.9 addresses directly.

2.2 Retrieval-Augmented Generation and Hybrid Retrieval

Retrieval-Augmented Generation (RAG) grounds LLM outputs in externally retrieved evidence rather than relying solely on parametric knowledge, and has become a primary strategy for improving factual reliability in generative systems (Lewis et al., 2020). Sawarkar et al. (2024) demonstrate that blending semantic search with hybrid query-based retrievers improves retrieval accuracy over either method alone, consistent with our design motivation for combining BM25 and dense retrieval where alphanumeric course codes require lexical precision that dense retrieval alone does not reliably provide. Beyond linear score combination, Cormack et al. (2009) introduce Reciprocal Rank Fusion (RRF), a rank-based (rather than score-based) fusion rule that avoids the cross-query scale-comparability problems of linearly combining two differently-distributed similarity scores; Section 3.5 uses RRF specifically for the query type (entity-heavy lookups) where a small number of high-confidence exact matches should dominate the ranking regardless of how the remaining candidates’ raw scores happen to be distributed for that particular query.

A growing literature addresses hallucination in RAG-based systems specifically. Huang et al. (2023) provide a taxonomy of hallucination causes and mitigation strategies, situating RAG as one of the most widely adopted inference-time mitigations. Varshney et al. (2023) propose validating low-confidence generations before they reach the user, closely related to this system’s confidence-gated abstention (Section 3.7). Joren et al. (2024) show that combining a context-*sufficiency* signal with retrieval/model confidence outperforms confidence alone on selective accuracy-coverage trade-offs, since confidence heuristics such as a top-1/top-2 score margin have blind spots – e.g., several equally-relevant near-duplicate candidates can produce a small margin despite unambiguous, correct retrieval. Section 3.7 implements exactly this combination using two sufficiency signals this system already computes as a side effect of retrieval (an exact metadata match and a graph-verified prerequisite chain), so no additional model call is required to obtain them. Es et al. (2023) propose RAGAS, a framework for evaluating RAG systems along axes (including faithfulness) that similarity-to-reference metrics cannot capture, since a faithful paraphrase and a confidently-stated fabrication can score identically on BLEU or ROUGE; Section 4.8 adopts RAGAS’s faithfulness formulation directly.

2.3 Domain Adaptation of Retrieval Embeddings

Reimers and Gurevych (2019) introduce Sentence-BERT, the siamese/triplet-network training scheme underlying the sentence-embedding models used throughout this system, and Henderson et al. (2017) introduce in-batch-negative contrastive training (the loss this paper’s embedding fine-tuning uses, via its `MultipleNegativesRankingLoss` instantiation), originally for response-suggestion retrieval. Off-the-shelf sentence embedding models are trained on general-domain text and are not specifically adapted to a given corpus’s own phrasing, entity distribution, or (in a bilingual corpus) cross-lingual question/answer pairing; Section 3.4 fine-tunes directly on this corpus’s own question/answer pairs, including its Banglish-question/English-answer pairs, and reports the held-out accuracy this produces before evaluating it inside the full pipeline.

2.4 Code-Mixed and Bilingual NLP

Code-mixed (or code-switched) text, where a single utterance mixes two languages or a transliterated form of one, is well documented as a distinct challenge for NLP systems trained primarily on monolingual data (Khanuja et al., 2020). For information retrieval specifically, a natural response is dictionary-based spelling-variant normalization before indexing and querying, so that lexical (term-frequency) matching is not defeated by inconsistent transliteration of the same underlying word. Section 4.9 implements and directly ablates this technique for our BM25 stream, and reports that it does not produce a significant effect on this particular test set, together with a structural explanation for why.

Closest to our own setting, Kruk et al. (2025) present BanglAssist, a Bengali-English code-switching customer-service chatbot, and report that “bilingual Banglish RAG” is not, however, an unclaimed research space – so we differentiate our contribution explicitly rather than assume novelty by domain alone. BanglAssist retrieves with a vector-only index reranked by a generic cross-encoder and translates every query to English before retrieval; our system retrieves with true hybrid BM25+dense fusion directly over Banglish text, motivated by exactly the reranker and translation ablations in Sections 4.4 and 3.12 that found a generic reranker and translation-before-retrieval do not help on our corpus. BanglAssist generates with a proprietary cloud API (GPT-4o); our system runs entirely on a local, open-weight 8B model, evaluated at 200 English and 100 Banglish queries with paired significance testing throughout, versus BanglAssist’s 20-query evaluation (6 Bengali, 9 Banglish, 5 English) with qualitative scoring and no significance testing. Most substantively, BanglAssist’s own description contains no entity-disambiguation or confidence-calibrated abstention mechanism – its only confidence-adjacent component is a fixed 0.8 cosine-similarity threshold gating direct FAQ-style matches, with no handling for an ambiguous entity reference or a calibrated decline. This is precisely the gap our ambiguous-entity notice and `conditioning.v2` mechanism (Section 3.8), together with calibrated confidence-gated abstention (Section 3.7), are built to close, in an academic-advising domain whose structured entities (course codes, prerequisite chains, faculty records) BanglAssist’s customer-support FAQ setting does not require reasoning over at all.

3 SYSTEM ARCHITECTURE AND METHODOLOGY

3.1 Dataset Description

The knowledge base comprises 5,059 records across seven tables: 2,297 EnglishQA question/answer pairs, 1,053 BanglishQA pairs (a Banglish question paired with its English-language answer, since the answer content itself is not duplicated per language), 814 FacultyAvailability records, 586 CourseDetails section records (course, section, faculty, class time, and room), 223 FacultyList records (name, initial, designation, office room, and email per faculty member), 52 Coordinator records, and 34 Prerequisite records, spanning topics from admission through thesis and internship stages. Figure 1 shows this distribution. EnglishQA and BanglishQA are each split into train/validation/test partitions (EnglishQA: 1,823/215/259; BanglishQA: 835/100/118) at collection time; Section 3.4’s embedding fine-tuning uses *only* the train partition, and every evaluation reported in Section 4 uses *only* the test partition, so no query scored in this paper was ever seen during fine-tuning. FacultyList was, in an earlier pass, ingested and embedded but never reachable via the exact-match mechanism (Section 3.3) that every other structured table relies on for entity-heavy precision – Section 4.5 reports the fix and its effect directly. FacultyAvailability remains ingested and retrievable but is not the source of any query in the automated-metric test sets scored in this paper; Section 4.6 reports a small dedicated evaluation against it instead.




Figure 1: Corpus composition across its seven source tables.

3.2 Data Ingestion Pipeline

The ingestion pipeline extracts data from CSV/spreadsheet sources into a relational (SQLite) store for structured access and auditability, and separately embeds it into ChromaDB for semantic search. Each vector record stores a text chunk, a unique identifier, and metadata (table of origin, course code, section, and question text where applicable) to support hybrid filtering and source attribution. A SHA-256 content hash is compared against stored records so that unchanged content is skipped rather than re-embedded on every ingestion run, keeping routine updates fast (Section 4.11 reports the resulting timings).




Figure 2: Workflow of database population, from source ingestion through change detection to vector storage.

3.3 Base Hybrid Retrieval

Figure 3 summarizes the two-stream retrieval workflow this section describes, which every component in Sections 3.4–3.7 builds on rather than replaces.

Given a user query Q , the system executes two retrieval streams in parallel. The lexical stream applies the Okapi BM25 ranking function (Robertson and Zaragoza, 2009) over the tokenized corpus. The semantic stream embeds Q with a SentenceTransformer model (Section 3.4) and retrieves the top-10 candidates by vector proximity in ChromaDB. Query tokens matching a course-code pattern are additionally looked up directly against each record’s course-code metadata (including a section-specific pattern, e.g. CSE111-07, a small alias table mapping informal course names to canonical codes, and – as of this revision – a faculty index matching initials and full names against the FacultyList table, see Section 4.5); an unambiguous exact match – exactly one candidate record matches the extracted code or identity – is treated as authoritative structured-database evidence rather than a probabilistic ranking signal and is guaranteed




Figure 3: Base retriever workflow: parallel BM25 and vector retrieval streams merged by weighted score fusion; Section 3.5 routes each query to one of two fusion configurations of this same mechanism.

rank-1 in its stream, on top of the additive exact-match bonus that already applied when more than one candidate shares the match. When a query names a bare course code together with a table-specific keyword (“prerequisite”, “chain”, “coordinator”), the match is additionally routed to that single table’s row rather than every table sharing the code, for the same reason: a query asking specifically about a course’s prerequisite is asking about one row, not the $\sim 10\text{--}30$ CourseDetails section rows that happen to share its code. Section 4.5 reports a direct, isolated test of both mechanisms’ marginal contribution, including a real coverage gap in the original version of this mechanism that the test itself surfaced.

BM25 and vector scores are normalized and combined via a convex combination, $\text{Score}(D, Q) = \lambda \cdot S_{\text{BM25}}(D, Q) + (1 - \lambda) \cdot S_{\text{vec}}(D, Q)$, exactly as in the base system this paper extends; $\lambda = 0.5$ (full hybrid), $\lambda = 1$ (BM25-only), and $\lambda = 0$ (vector-only) remain the three configurations directly compared in the ablation (Section 4.1). An additional rank-based fusion mode, Reciprocal Rank Fusion (Cormack et al., 2009), is available and is what the adaptive router (Section 3.5) selects for entity-heavy queries.

3.4 Domain-Adapted Bilingual Embeddings

The base sentence-embedding model (all-MiniLM-L6-v2) is general-purpose and was never trained on this corpus’s own phrasing, course-code-heavy content, or its Banglish-question/English-answer pairing. We fine-tune it via MultipleNegativesRankingLoss (Henderson et al., 2017), an in-batch-negatives contrastive objective, using every

Table 1: Held-out validation retrieval accuracy by embedding model ($n = 540$ question/answer pairs, Split=val, never trained on; single consistent run, re-measured 2026-08-01 after BanglishQA grew from 1,053 to 3,044 rows).

Model	Top-1	Top-5	MRR
Base (pretrained, not fine-tuned)	0.446	0.819	0.605
Fine-tuned MiniLM, in-batch negatives only (previously deployed)	0.513	0.920	0.677
Fine-tuned MiniLM, mined hard negatives, pre-Banglish-expansion (previously deployed)	0.524	0.909	0.682
Fine-tuned MiniLM, mined hard negatives, Banglish-expanded (deployed)	0.520	0.928	0.689
Fine-tuned BanglishBERT	0.213	0.557	0.367
Fine-tuned multilingual-MiniLM	0.522	0.874	0.669

(question, answer) pair from the EnglishQA and BanglishQA *train* splits (2,658 pairs) as positive pairs, 3 epochs, batch size 32. Validation and test splits are never used for training.

To check that this domain adaptation is a genuine improvement and not merely a plausible design choice, we evaluate question-to-answer retrieval accuracy on the held-out *validation* split (540 pairs as of this revision, never trained on) before deploying the model: for each validation question, we rank all validation answers by cosine similarity and check whether the question’s own answer is ranked first (Top-1), within the top 5 (Top-5), or compute its mean reciprocal rank (MRR). Table 1 reports this comparison against the un-fine-tuned base model and against two alternative fine-tuning targets: a Bengali-transliteration-aware pretrained model (BanglishBERT) and a multilingual sentence model (paraphrase-multilingual-MiniLM-L12-v2), each fine-tuned identically on the same train split, alongside two successive mined-hard-negatives checkpoints described below. All rows in this revision of the table were re-measured together in one consistent execution of `scripts/eval_embeddings_held_out.py` (confirmed exactly reproducible by running the base-model row twice independently and diffing the output); they supersede a subset of numbers from an earlier, separately-run pass whose exact code state at the time is no longer reconstructable and which we no longer treat as authoritative. This re-measurement (2026-08-01) also caught a real gap: the currently-deployed checkpoint was missing from the evaluation script’s own model list entirely, the same class of omission already found and fixed for the structured-table held-out check (Section 3.4.1) – added here rather than left unnoticed a second time. The $n = 297 \rightarrow 540$ growth itself is not a measurement change but a genuine corpus change: BanglishQA nearly tripled (1,053 to 3,044 rows) in a data-ingestion pass after the original table was built, which is also why the currently-deployed checkpoint (trained on this expanded set, see below) did not exist as a row in the original table.

Fine-tuning `all-MiniLM-L6-v2` on this corpus’s own pairs remains the strongest family of configurations by a clear margin over the two alternative base models, and some variant of it is what Sections 4.1–4.9 deploy. Within that family, the currently-deployed Banglish-expanded checkpoint is not the single best on Top-1 (0.520, essentially tied with the pre-expansion checkpoint’s 0.524 and the multilingual model’s 0.522) but is best on Top-5 and MRR, the two rank-tolerant metrics – consistent with it being trained on a validation set now much more heavily weighted toward Banglish content (317 of 540 pairs) than when the pre-expansion checkpoint was deployed, and we report this honestly rather than claim a clean sweep that the current numbers do not show. Both alternative fine-tuning targets still underperform the MiniLM family: BanglishBERT, despite being pretrained specifically on Bengali-transliterated text, scores lowest of all six configurations after fine-tuning on this same data, and the multilingual model remains only marginally above the un-fine-tuned English-only base on Top-1 despite being competitive on Top-5. We report this as a genuine negative result for the intuitive hypothesis that a Bengali-aware or multilingual base model should out-perform an English-only one on a bilingual corpus – on this data, it does not, and corpus-specific fine-tuning of the simpler English-only model matters more than the base model’s language coverage.

3.4.1 Mined Hard Negatives

The originally deployed fine-tuning (`MultipleNegativesRankingLoss` with only (anchor, positive) pairs) draws its only negative signal from whatever unrelated answers happen to land in the same training batch – almost always trivially easy to distinguish from the true answer once a batch is even moderately diverse, so the model is rarely forced to learn the fine-grained distinctions this corpus specifically needs (near-duplicate paraphrase rows for the same fact; structurally similar but factually distinct rows, e.g. different course codes’ prerequisite answers or different faculty members’ schedule answers, that share most of their surface vocabulary), following the general finding of Meghwani et al. (2025) that in-batch random negatives under-train exactly this kind of fine-grained discrimination. We mined one explicit hard negative per training anchor: the single most cosine-similar *other* answer under the un-fine-tuned base model that is not a near-duplicate of the true answer (excluded by an exact-text check, so a hard negative is never actually another correct paraphrase of the same fact), then trained with the same `MultipleNegativesRankingLoss` extended with this explicit negative alongside the usual in-batch ones – strictly more negative signal per training step, not a replacement of the original procedure.

This is a genuine, statistically confirmed improvement: mined hard negatives beat the previously-deployed in-

batch-only fine-tuning on both MRR (0.536 vs. 0.502, paired bootstrap mean diff +0.034, 95% CI [0.006, 0.063], $p = 0.018$) and Top-1 (0.290 vs. 0.246, mean diff +0.044, 95% CI [0.003, 0.084], $p = 0.030$), on the 297 held-out pairs available at that time, paired at the level of each pair’s own rank. Having validated this at the embedding level, we then deployed it: the vector corpus (5,059 chunks) was re-embedded and the ChromaDB index rebuilt under the new model, and every retrieval-only evaluation in this paper (Sections 4.2, 4.7.2, 4.7.3, 4.7.1) was re-run against the rebuilt index to confirm nothing regressed – reported inline in those sections rather than assumed. A stratified $n = 20$ generation-quality sample (Section 4.1) was also re-run as a smaller, honestly-scoped confirmatory check, since the full 800-generation regeneration was not time-feasible this revision.

A second retrain followed a subsequent data-ingestion pass that nearly tripled BanglishQA (1,053 to 3,044 rows): fine-tuning an off-the-shelf sentence-embedding model on additional code-mixed text is not a guaranteed win in general (more training data of a harder, noisier kind can degrade performance rather than improve it, depending on the base model and data quality), so we did not assume more Banglish training data would help and instead re-fine-tuned on the expanded set and tested the result directly, on three held-out sets, against the checkpoint it would replace: a paired bootstrap comparison (2,000 resamples) finds the Banglish-expanded checkpoint significantly better on the Banglish-only held-out set (MRR +0.026, 95% CI [0.003, 0.049], $p = 0.023$; Top-1 +0.035, 95% CI [0.003, 0.069], $p = 0.031$, $n = 317$), the direct question this retrain targets, with no significant regression on the pooled QA-pair set (MRR $p = 0.568$, Top-1 $p = 0.819$, $n = 540$) or the structured-table set (both checkpoints score a perfect 1.000/1.000, $n = 328$). This is the outcome we hoped for but did not assume: a real, targeted gain with no measured collateral cost, so this checkpoint (`finetuned_minilm_hard_negatives_structured_banglish_expanded`) is what is actually deployed as of this revision – Table 1 reports it directly rather than the intermediate pre-expansion checkpoint. Every number in this paper from here onward reflects this deployed embedding model unless explicitly marked otherwise.

3.5 Adaptive Fusion Routing

The base system’s own λ -sensitivity sweep (reported in the system’s earlier evaluation and reproduced in spirit by Section 4.7 below) found no single fixed λ that significantly beats BM25-only on either entity-heavy or open-ended queries, on any of four automated metrics, across the full swept range. We treat this as evidence that a single global fusion rule is the wrong model for this corpus, not that lexical retrieval should simply always win, and instead route each query to one of two fusion configurations by the same entity-heavy/open-ended signal the sweep was already stratified by (an entity-heavy query is one containing a course-code token or a known alias): entity-heavy queries are routed to linear fusion at $\lambda = 0.9$ (heavily lexical-weighted), and open-ended queries are routed to the original, already-validated linear fusion at $\lambda = 0.5$. Entity-heavy queries were originally routed to Reciprocal Rank Fusion at the same $\lambda = 0.9$ instead; Section 4.7.1 reports why that choice was replaced.

3.6 Prerequisite-Graph Augmentation

Multi-hop prerequisite questions (“what do I need before I can take CSE310?”) require chaining several one-hop prerequisite edges; a 5-hop chain, for instance, cannot be recovered by retrieving the single chunk that states CSE310’s immediate prerequisite, nor does a short text chunk have room to state a full transitive chain even if one existed verbatim in the corpus. We build a directed graph from the Prerequisites table’s own one-hop (course, prerequisite) edges via `networkx`, rather than from the table’s separate free-text “full chain” column, which we found to be inconsistent with the one-hop edges at the data-entry level for some courses. For a query recognized as asking about a full prerequisite *chain* (as opposed to a single course’s direct prerequisite, which plain retrieval already answers correctly, since it is exactly what one corpus chunk states), the graph’s breadth-first transitive closure is computed and prepended to the retrieved context as a verified chain, before generation.

Only 12 of the 200 English test queries actually trigger this component (verified directly by calling the graph module against every test query), so we isolated it: the full pipeline, with `use_graph` on vs. off, run on exactly those 12 queries, everything else held fixed. A first run of this isolated ablation found graph augmentation *significantly hurt* BLEU ($p = 0.025$, paired t -test) – surprising, since the graph provides strictly correct information the alternative (plain retrieval) frequently could not. Inspecting individual outputs found the cause immediately: the graph block stated chains with arrow notation (“CSE370 → CSE221 → ...”), while this corpus’s own reference answers and source `FullChainPreRequisite` field use hyphens (“CSE370-CSE221-...”); the model would adopt the graph block’s arrow formatting in its answer, and BLEU, an n -gram-overlap metric, credits none of the overlap between the two equally-correct notations. We changed the graph block to hyphen-join its chain, matching the corpus’s own convention, and re-ran the isolated ablation from scratch.

The main test set’s 12 trigger queries are a real but small subset of what the corpus actually supports: 29 courses have a genuine, non-empty, database-verified full prerequisite chain, computed by the same graph-traversal code the deployed system uses, not a hand-typed approximation. We grew the isolated ablation to $n = 31$ by adding the 19 chains not already exercised by the main test set – real facts already present in the knowledge base, not synthetic or fabricated data – and re-ran the identical on/off comparison.

Table 2: Prerequisite-graph augmentation, isolated ablation. Rows track the ablation’s history: the arrow-notation bug (pre-fix), the hyphen-notation fix at the original $n = 12$ trigger set, and the current, expanded $n = 31$ measurement (12 original plus 19 additional real, database-verified chains not previously tested) under the current retriever.

Condition	BLEU	ROUGE-L	BERTSc.	METEOR
Graph on (arrow notation, pre-fix, $n=12$)	0.599	0.928	0.967	0.869
Graph on (hyphen notation, $n=12$)	0.917	0.961	0.992	0.960
Graph off (plain retrieval, $n=12$)	0.755	0.869	0.968	0.877
Graph on (current, $n=31$)	0.941	0.974	0.994	0.977
Graph off (current, $n=31$)	0.730	0.859	0.965	0.869

Table 3: Calibrated per-route abstention thresholds (against labeled Out-of-Scope vs. answerable examples, seed=42; entity-heavy augmented with 52 database-verified synthetic examples, see text).

Route	Signal	Thr.	Prec.	Rec.	n
Entity-heavy	margin	99.99	0.900	1.000	60
Open-ended	top-1 score	1.028	0.657	0.870	372

After the notation fix, graph augmentation’s point estimate was higher than plain retrieval on all four metrics at $n = 12$, reversing the pre-fix BLEU regression entirely, but re-measurement under the current retriever (exact-match fix, BM25 retuning) left the significance picture mixed at that sample size: BLEU significant under a paired t -test ($p = 0.042$) but not Wilcoxon ($p = 0.063$); BERTScore the reverse ($p = 0.031$ Wilcoxon, $p = 0.078$ paired- t); ROUGE-L and METEOR non-significant under both. At $n = 12$ this did not clear this paper’s own both-tests-agree bar.

The expanded $n = 31$ measurement clears it cleanly, on every metric: BLEU (+0.211, $p < 0.0001$ both tests), ROUGE-L (+0.116, $p = 0.0009/p < 0.0001$), BERTScore (+0.029, $p = 0.0004/p < 0.0001$), and METEOR (+0.108, $p < 0.0001$ both tests). This is a genuine, confirmed win, not a re-measurement of the same 12 queries under friendlier conditions – 19 of the 31 queries are facts the main test set never exercised before, and the effect replicates cleanly across the enlarged set. We report both the bug and the fix, and both the underpowered and the expanded measurement, rather than only the most favorable number: the pre-fix result would otherwise have been read as genuine evidence against graph augmentation, when it was in fact a formatting artifact, and the $n = 12$ result on its own would have understated a real effect that a larger, still entirely real, sample confirms.

3.7 Confidence-Gated Abstention with Sufficient-Context Override

Every retrieval call already computes a margin between its top-1 and top-2 fused scores, and the raw top-1 score, as candidate confidence signals. We calibrate a decline-to-answer threshold per route (entity-heavy vs. open-ended) against a labeled set drawn from the corpus’s own “Out of Scope / Unanswerable” category, whose reference answers are deliberate refusals rather than facts, split against an equal-sized sample of answerable queries (seed fixed at 42). A single global threshold was found, in an earlier pass, to systematically over-abstain on the entity-heavy/RRF route and under-abstain on open-ended/linear, since RRF’s top-1/top-2 margins are structurally an order of magnitude smaller than linear fusion’s even at equal retrieval confidence; per-route calibration corrects this.

The entity-heavy route’s natural calibration sample is small: only 18 of the 380 Out-of-Scope/answerable rows happen to mention a course code or faculty identity at all, most Out-of-Scope rows being generic questions unrelated to a specific entity. Rather than accept a wide confidence interval as final, we grew this sample with additional, non-fabricated labeled examples: a query naming a course code verified (by direct database lookup) to not exist anywhere in CourseDetails/Prerequisites/Coordinator is unambiguously should-abstain by construction, and a query naming a real code that does have Prerequisites/Coordinator data is unambiguously answerable; the same logic applies to faculty initials against FacultyList. This added 52 template-generated (not organically collected) examples, disclosed here as such. Table 3 reports the resulting calibration, after this augmentation and after the exact-match/routing fix of Section 4.5.

Growing the entity-heavy calibration set from 18 to 60 examples tightens its 95% accuracy CI from (0.65, 0.99) to (0.861, 0.990) – still not large, but a substantially more precise estimate of a threshold that, mechanically, is now set at the unambiguous-match ceiling itself (Section 4.5): once every entity-heavy query produces a real margin dominated by that ceiling, the confidence signal and the structural exact-match signal are no longer independent, which we note as a modeling simplification rather than a problem this table hides. The open-ended route’s recall improved from 0.457 (pre-fix retriever) to 0.870 after recalibrating against the fixed retriever’s confidence distribution – the confidence-threshold gate alone now catches most, not fewer than half, of genuinely out-of-scope questions on the calibration set.

Following Joren et al. (2024), we do not rely on the confidence threshold alone: this system already computes two

structural sufficiency signals as a side effect of retrieval, at no extra cost – whether any chunk in the assembled context is an exact metadata match, and whether the prerequisite graph (Section 3.6) produced a verified chain – and either one overrides a would-be abstain decision, on the grounds that direct structural evidence the corpus contains the answer should outrank a confidence heuristic already known to have blind spots on this corpus (e.g., several equally-good near-duplicate paraphrase chunks producing a small top-1/top-2 margin despite unambiguous, correct retrieval). This override is applied only over chunks actually assembled into the context the generator receives, not merely present at some earlier retrieval rank, so an abstain decision is never overridden on evidence the model was never shown.

A code review after this fix found the entity-heavy abstention gate had become more brittle in one specific respect: since the unambiguous-match ceiling (Section 4.5) now dominates the confidence margin whenever it fires, the entity-heavy route’s abstention decision has effectively collapsed into a binary “did exact-match fire” signal, and any answerable entity-heavy query that does not trigger exact-match – because of how the entity happened to be written, not because the corpus lacks the answer – is at risk of always being declined. Directly testing this, we found exactly such a case: “CSE 220” and “CSE-220” (with a space or dash) matched no course code at all under the original pattern, only the glued form “CSE220” did, confirmed live before any fix. We closed this specific gap with a small regex extension (an optional space/dash separator, Section 3.10), verified to restore exact-match on both forms without any other behavior change. This does not fully resolve the underlying brittleness – a misspelled faculty name or an unanticipated formatting variant regex cannot enumerate would still fall through – so we also implemented, but did not enable by default, an LLM-based entity-normalization retrieval fallback motivated directly by Magomere et al. (2025)’s claim-normalization mitigation for embedding robustness against input perturbation: for a query where exact-match finds nothing but a loose entity-shaped pattern is still present, the query is rewritten to a canonical form by the same local LLM before a second retrieval pass, unioned with the original (never replacing it, the same safe pattern as query translation, Section 3.12). This closes the general class of the problem in principle; we report it as implemented-but-unvalidated rather than claim a benefit not yet measured by its own ablation.

3.8 Ambiguous-Entity Disambiguation

Course-code lookups in this corpus are collision-free by construction (verified directly: 0 collisions among 586 `CourseDetails` rows across 71 distinct base course codes), but faculty-name lookups are not – a query naming a common surname (e.g. “Rahman”) can match several distinct `FacultyList` records. The base exact-match mechanism (Section 4.5) is not designed to resolve this: when more than one candidate matches, no single record is authoritative, and a generator given all of them in context with no further instruction risks confidently naming just one, silently arbitrary among several equally-valid candidates – worse for an advising chatbot than an answer that visibly hedges, since a wrong confident answer is indistinguishable from a right one to the student reading it.

We address this with two additive context blocks, injected at the same `float("inf")` confidence priority as the prerequisite-graph and faculty-room blocks (Sections 3.6, 4.6) whenever more than one candidate resolves to the same queried entity: a flat notice instructing the model not to silently pick one candidate and to ask the user to disambiguate instead, and a conditioning hint that checks whether the tied candidates differ on a distinguishing field this corpus already stores (`Designation`, `Initial`, `Status`, checked in that priority order) and, if so, surfaces it directly in a concrete clarifying-question template (e.g. “Which one do you mean: Professor or Lecturer?”) rather than requiring the user to re-type a full name.

An initial version of the conditioning hint phrased this as a conditional instruction (“if the user’s phrasing already narrows it down by *field*, use that”). Evaluated against a flat-notice-only baseline and a no-notice control on 55 genuinely-ambiguous faculty-name queries with an LLM judge (three criteria: does the response avoid stating one answer with false confidence; does it ask the user to clarify; does it offer a distinguishing attribute), both the flat notice and this original conditioning hint significantly improved over no-notice on the core safety criteria, but the conditioning hint’s own specific selling point – offering a distinguisher – did not reach significance over the flat notice alone ($n = 55$). Expanding the test set to $n = 220$ (the same 55 ambiguous people, four distinct field questions each, so more real distinct questions rather than repeated padding) confirmed this was not a power issue: the original conditioning hint remained statistically indistinguishable from the plain flat notice on all three criteria.

Rather than treat this as the final word, we redesigned the hint to be more directive – giving the model the distinguishing values pre-assembled into a ready-to-use question rather than a conditional instruction to check for them – and re-tested this second version (“conditioning_v2”) against all three prior conditions on the same $n = 220$ set. A methodological check on the judge itself preceded trusting the result: the first significance pass used the original, holistic `avoids_false_confidence` judge (already known unreliable at chance-level agreement with human labels on a subset check), and it produced a striking but suspicious pattern – `conditioning_v2` scored significantly *worse* on `avoids_false_confidence` while significantly better on the other two criteria. Manual inspection of the flagged low-scoring responses found textbook-correct clarifying questions (e.g. “Which one do you mean: C. Lecturer (Contractual) or Lecturer (Full Time)?”) being scored as failing to avoid false confidence, backwards relative to the criterion’s own definition. Re-scoring `avoids_false_confidence` with the validated, decomposed judge already built for the main faithfulness evaluation (Section 4.8) changed the verdict on 43.3% of responses, reversing the apparent regression

Table 4: Ambiguous-entity notice quality, LLM-judge criteria (0–1, higher is better), $n = 220$ expanded faculty-name-collision test set, paired bootstrap significance.

Condition	Avoids false conf.	Asks to clarify	Offers distinguisher
No notice	0.486	0.027	0.118
Flat notice	0.805	0.586	0.327
Conditioning (v1)	0.818	0.573	0.332
Conditioning_v2 (deployed)	0.864	0.673	0.509

entirely.

Table 4 reports the corrected result. `conditioning_v2` significantly outperforms both the flat notice and the original conditioning hint on all three criteria ($p < 0.0001$ throughout, paired bootstrap, $n = 220$). This is deployed as of this revision (`pipeline/novel_pipeline.py`’s `build_conditioning_hint_v2`) – the original conditioning hint is superseded at the call site but not deleted, so the comparison in Table 4 remains exactly reproducible.

This mechanism is, to our knowledge, a genuine differentiator against the closest related system: BanglAssist (Section 2) has no entity-disambiguation or ambiguity-aware confidence mechanism at all, only a fixed cosine-similarity threshold gating direct FAQ matches. We note two honest limitations. First, this test set is faculty-name collisions only; course-code collisions do not currently exist in this corpus to test the same mechanism against (Section 4.5), so the result is validated on one collision type, not proven to generalize to a structurally different one. Second, the judge-reliability correction above is itself evidence that automated LLM-judge scoring of nuanced, definition-sensitive criteria like “avoids false confidence” is not safe to trust without a validated, decomposed judge or a cross-check – we disclose this as a limitation of the evaluation methodology broadly, not specific to this one result, consistent with Section 4.8’s independent NLI cross-check for the same reason.

3.9 Confidence-Ordered Context Assembly

A second, independent risk affects every configuration that assembles context from more than one source (the prerequisite graph, sufficient-context pool matches, and the ranked retrieval results themselves): Liu et al. (2024) show that language models use information placed at the start or end of a long context far more reliably than information buried in the middle, regardless of whether it technically fits within the context window. Prior to this revision, context was assembled in a fixed order – graph block, then pool matches, then ranked results – irrespective of confidence, so the single highest-confidence retrieved chunk routinely ended up in the middle of the assembled prompt whenever a graph block or pool match also existed, exactly the position this literature identifies as least reliable.

We also confirmed the context-length side of this risk directly: `num_ctx` (Section 3.13) was set to 2048 against an earlier measurement of the maximum retrieved-context length (2508 characters, taken before the prerequisite-graph and pool-match extensions existed); re-measuring against current outputs found the actual maximum has grown to 5,041 characters ($\approx 1,260$ tokens), with the 99th percentile at $\approx 1,140$ tokens – close enough to the old budget, once the system prompt and generated answer are added, to risk silent truncation with no detection mechanism in place. We raised `num_ctx` to 4096.

Raising the ceiling alone does not address the ordering risk, so we additionally implemented confidence-ordered context assembly: every piece of context (graph block, pool matches, ranked retrieval results) is assigned a confidence score – verified structural evidence (a graph block or an already-validated sufficient-context pool match) scores above any ordinary retrieval score, consistent with how the unambiguous-match ceiling already treats structural evidence as authoritative rather than merely probable – and the pieces are arranged in a zig-zag order: highest confidence first, second-highest last, third-highest second, working inward, following the retrieval-reordering approach of Jin et al. (2024). This directly reuses the same confidence signal the abstention gate already computes rather than introducing a second, unrelated one, and we see it as a small unifying design point across three previously separate concerns (Sections 3.7, 4.5, and this one): a single confidence score per chunk, computed once, that can gate abstention, order context, and – as a further extension we identify but do not implement in this revision – decide which low-confidence chunks are worth extractively compressing before insertion rather than included in full or dropped entirely, in the spirit of Hwang et al. (2025)’s sentence-level extractive compression for RAG. We describe this as a modest, honestly-scoped unification of existing techniques (graduated confidence scoring, zig-zag reordering, and structural-evidence prioritization are each separately published) applied consistently to this system’s own context-assembly step, not as a novel scoring or reordering algorithm in its own right.

3.10 Shared Patterns and a Table-Coverage Audit

Two of this revision’s fixes (Sections 4.5 and 4.6) shared a root cause worth naming directly: a single real-world concept (a course code; which table a faculty-related query actually wants) was represented by more than one independently-

Table 5: English retrieval ablation, single-method baselines, after the exact-match coverage fix of Section 4.5 ($n = 200$).

Variant	BLEU	ROUGE-L	BERTSc.	METEOR
Full Hybrid ($\lambda = 0.5$)	0.694	0.852	0.969	0.873
BM25-only ($\lambda = 1$)	0.655	0.820	0.965	0.859
Vector-only ($\lambda = 0$)	0.495	0.627	0.932	0.663
No-retrieval	0.046	0.191	0.865	0.290

maintained piece of code, and the two copies silently drifted apart. We addressed the specific instance `pipeline/prerequisite_graph` had its own copy of the course-code pattern that never received the letter-suffix fix applied to the retriever’s copy – by extracting a single shared pattern module (`pipeline/patterns.py`) that both now import, with a regression test asserting the two modules reference the literal same compiled pattern object, not merely equivalent behavior. We also built a standing table-coverage audit: for every table in the corpus, does at least one query across all existing test sets retrieve a chunk from that table at rank 1? Both of this revision’s table-routing bugs were, in retrospect, tables with zero such queries in any test set before their dedicated evaluations were built; running this audit now, after both fixes, finds no remaining zero-coverage table, but its value is as a standing check against the next one, not a one-time result.

3.11 Cross-Encoder Reranking (Tested, Disabled by Default)

A cross-encoder reranker (base model `cross-encoder/ms-marco-MiniLM-L-6-v2`) was implemented and tested as a candidate fifth component: after adaptive routing produces a pool of 10 candidates, the reranker jointly scores each (query, candidate) pair and re-sorts the top 5 before generation. Section 4.4 reports that this component, evaluated after Section 3.4’s embedding fine-tuning was already in place, is a significant net *negative*: the generic, non-fine-tuned reranker’s joint relevance judgments disagree with – and on this corpus, are measurably worse than – the already domain-adapted ranking the fine-tuned embeddings alone produce. We additionally fine-tuned this reranker on the corpus’s own hard negatives (real retrieved near-misses, contrastive binary-cross-entropy loss) to test whether domain adaptation would flip this result; Section 4.4 reports that it does not – the fine-tuned reranker remains a net negative, only marginally less so than the generic one. Both variants are implemented and available but disabled by default; every other result in this paper uses the pipeline with reranking off.

3.12 Cross-Lingual Query Translation and BM25 Code-Mixed Normalization

Two further components were implemented specifically for the Banglish setting, motivated by the cross-lingual RAG and code-mixed IR literature (Section 2), and both are reported honestly in Section 4.9 as producing no significant effect on this test set, together with an explanation for why. First, query translation: for queries a cheap keyword-based gate flags as likely Banglish, the query is translated to English by the same generation LLM before a second retrieval pass, whose results are unioned with the original-language retrieval (never replacing it, only adding candidates). Second, BM25-level normalization: a small dictionary collapses known Banglish spelling variants of the same root word (e.g. “korte”/“korbo”/“korch” → “kor”) to one canonical token before both indexing and querying, the same principle as English stopword removal applied to spelling variation instead of function words.

3.13 Generative Reasoning Engine

Response generation uses a locally hosted Llama-3.1-8B model served via Ollama. Retrieved chunks (and, where applicable, the prerequisite-chain block) are concatenated into a structured system prompt that separates retrieved context from the user query, directing the model to answer from the provided evidence. If the abstention gate (Section 3.7) fires, generation is skipped entirely and a fixed decline-to-answer message is returned instead.

4 EXPERIMENTAL EVALUATION

All automated-metric results below use BLEU (Papineni et al., 2002), ROUGE-L (Lin, 2004), BERTScore (Zhang et al., 2019), and METEOR (Banerjee and Lavie, 2005) scored against reference answers, with paired t -tests (and, where noted, matched Wilcoxon signed-rank tests) for significance testing. The English test set is 200 queries (100 open-ended, 100 entity-heavy) drawn from the test-split partitions described in Section 3.1; the Banglish test set is 100 queries drawn the same way from BanglishQA’s test split.



Figure 4: Workflow of response generation, from user input through adaptive hybrid retrieval, graph augmentation, and abstention gating, to the final LLM-generated answer.

4.1 English Retrieval Ablation

With the fine-tuned embeddings (Section 3.4) in place, the pattern differs from the base system’s original finding: full hybrid now leads BM25-only outright on all four metrics (Table 5) rather than being statistically tied with it, consistent with the mechanism identified in Section 3.4 – a stronger dense signal can contribute more to the blend than a weak, un-adapted one could. Vector-only remains clearly behind both, and no-retrieval remains far behind all three, exactly as before. All four configurations improve marginally over an earlier measurement round following the exact-match coverage fix (Section 4.5), which affects any configuration with $\lambda > 0$.

Table 5 reflects the embedding model in place *before* Section 3.4.1’s hard-negative-mining deployment; regenerating it in full under the new embedding model (all 800 generations) was not feasible in this revision, at an estimated 5.5–10 hours of CPU-only inference time (`scripts/run_ablation.py`’s own documented estimate). As a smaller, honestly-labeled confirmatory check, we instead ran a stratified $n = 20$ subset (10 entity-heavy, 10 open-ended, fixed seed) under the redeployed model: Full Hybrid 0.585/0.740/0.952/0.770, BM25-only 0.488/0.697/0.947/0.746, Vector-only 0.371/0.501/0.913/0.530, No-retrieval 0.036/0.182/0.861/0.256 (BLEU/ROUGE-L/BERTScore/METEOR). Every configuration scores below its Table 5 value on this subset, including No-retrieval, which touches no embedding model or retrieval step at all – since an embedding-model change cannot possibly affect a configuration that never embeds anything, this uniform drop is attributable to this particular 20-query sample happening to be a harder-than-average subset of the 200, not to the new embedding model, and we flag it as exactly that rather than let it look like a regression. The relative ordering the paper’s other results depend on (Full Hybrid > BM25-only > Vector-only > No-retrieval) holds identically on all four metrics in this subset, which is the property we were actually checking for.

4.2 Standard IR Metrics: Recall@k, MRR, nDCG@5/@10

Beyond end-to-end generation quality, we separately report retrieval quality in the standard IR reporting format (Recall@1/3/5, MRR, nDCG@5 and nDCG@10), with binary relevance judged per query: for open-ended queries, a candidate is relevant iff the (verbatim, corpus-native) reference answer is a substring of its text; for entity-heavy queries, relevance is judged against the reference answer’s own course-code/email/faculty-identity entities (Section 4.5’s methodology, reused directly). We report nDCG@5 as the headline cutoff alongside nDCG@10, matching the convention of SemEval-2026 Task 8 (MTRAGEval), whose system-description papers (e.g., Athanasiou et al. (2026)) report nDCG@5 as the primary retrieval metric with nDCG@10/Recall@10 secondary. nDCG’s ideal ranking (IDCG) is computed from the actual number of relevant candidates found in each query’s own top- k , not assumed to be exactly one – this corpus’s substantial near-duplicate/paraphrase content means more than one retrieved chunk is often genuinely relevant, and assuming otherwise silently produces nDCG values that can exceed 1.0, an error we caught and corrected before reporting these numbers.

Paired bootstrap significance testing (2,000 resamples) between the adaptive router and each baseline shows a qualitatively similar picture to the generation-quality ablation, with one notable difference: adaptive is statistically indistinguishable from full hybrid on every metric, and from BM25-only on Recall@1/3/5 and MRR, but the generation-level ROUGE-L/BERTScore advantage adaptive shows over BM25-only (Section 4.3) is *not* confirmed at this retrieval-only

Table 6: Retrieval-only IR metrics, all four configurations ($n = 200$); adaptive uses the production `retrieve_adaptive` routing. Re-measured after deploying the mined-hard-negatives embedding model (Section 3.4.1) and rebuilding the vector index; see discussion below for what changed vs. the previous embedding model.

Config	R@1	R@3	R@5	MRR	nDCG@5	nDCG@10
Adaptive	0.910	0.920	0.920	0.915	0.878	0.899
Full Hybrid	0.910	0.920	0.920	0.915	0.884	0.902
BM25-only	0.905	0.915	0.920	0.911	0.870	0.895
Vector-only	0.560	0.600	0.605	0.585	0.574	0.596

Table 7: Adaptive pipeline (reranker off, deployed default) vs. baselines, matched pairs only ($n = 199/200$; 1 excluded where the pipeline abstained), re-measured 2026-08-01 under the current retriever (exact-match coverage fix and BM25 k_1/b retuning both applied) after the previous measurement (2026-07-31, predating both fixes) was found to be stale. Positive mean diff favors the adaptive pipeline; Wilcoxon p in parentheses where it disagrees with the paired t -test.

Comparison	Metric	Mean diff	p
vs. Full Hybrid	ROUGE-L	−0.017	0.205
vs. Full Hybrid	METEOR	−0.013	0.341
vs. BM25-only	BERTScore	−0.0035	0.097 (0.808)
vs. BM25-only	METEOR	−0.025	0.041 (0.306)
vs. Vector-only	BLEU	+0.012	0.480
vs. No-retrieval	BLEU	+0.610	< 0.001

level on either nDCG cutoff (nDCG@5 mean diff +0.009, 95% CI [−0.007, 0.024]; nDCG@10 mean diff +0.004, 95% CI [−0.003, 0.012]; neither significant) – suggesting that advantage arises somewhere in generation given equivalent-quality retrieved context, not from a detectably better ranking. Vector-only is significantly behind adaptive on every metric ($p < 0.001$), driven almost entirely by entity-heavy queries: broken out by query type, vector-only’s entity-heavy Recall@1 is 0.17 versus 0.85–0.97 for every other configuration, while its open-ended Recall@1 (0.95–0.97) is competitive with the rest – the sharpest possible illustration of this paper’s original design motivation, that dense retrieval alone is specifically weak on exact-identifier queries.

Comparing this table against its previous measurement (under the in-batch-negatives-only embedding model, Section 3.4.1) isolates exactly what the new embedding model changed at the full-pipeline level, beyond the isolated held-out QA-pair check reported earlier: BM25-only, Full Hybrid, and Adaptive are essentially unchanged (all differences ≤ 0.008 , well within the near-ceiling noise these configurations already showed, since all three are dominated by the exact-match mechanism, Section 4.7.2), while Vector-only shows a small, mixed shift – Recall@1 improves slightly (0.550 $\rightarrow$ 0.560) but Recall@5 and both nDCG cutoffs move slightly down (0.660 $\rightarrow$ 0.605; 0.586 $\rightarrow$ 0.574; 0.606 $\rightarrow$ 0.596). We report this honestly rather than only the isolated win: the held-out embedding-level evaluation (Table 1) trains and validates on (question, answer) *pairs* specifically, while this corpus’s actual vector index also contains many chunks derived from structured tables (course listings, faculty schedules) that are not in that distribution – the hard-negative mining likely sharpened the model’s discrimination among QA-pair-like text specifically, at a small cost to how it ranks the corpus’s non-QA-pair chunks beyond the top-1 slot. Since Vector-only is a diagnostic baseline, not a deployed configuration, and every deployed configuration (Full Hybrid, Adaptive) is unaffected or marginally improved, this does not change the deployment decision, but it is a genuine limitation of the isolated held-out check as a complete predictor of full-pipeline behavior, worth flagging for future embedding work on this corpus (e.g. mining hard negatives from structured-table chunks too, not only QA pairs). We in fact already did this: a later fine-tuning pass explicitly extended hard-negative mining to include structured-table pairs alongside QA pairs, then further to the Banglish-expanded corpus (`scripts/finetune_embeddings_banglish_expanded.py`), and this is the model actually deployed today (`pipeline/chroma_embedding.py`’s `DEFAULT_MODEL`). Re-running the structured-chunk held-out check (`scripts/eval_structured_embeddings.py`, $n = 328$) directly against this deployed model, rather than the earlier QA-pairs-only checkpoint the original version of this finding described, shows the regression does not apply to it: Top-1/Top-5/MRR of 1.000/1.000/1.000, identical to the structured-extended intermediate checkpoint and a full recovery from the QA-pairs-only model’s 0.466/0.585/0.524. We verified this directly rather than assume the later retrain also fixed it, since an embedding-model regression on exactly the structured-table content this system’s entity-heavy contributions depend on would otherwise be a real, undisclosed risk.

4.3 Adaptive Pipeline vs. Baselines

This table has itself been re-measured since the previous revision, for the same reason as the reranker rows above: the earlier version (2026-07-31) predated both the exact-match coverage fix (Section 4.5) and the BM25 k_1/b retuning

Table 8: Cross-encoder reranker on (fine-tuned) vs. off, otherwise identical pipeline, matched vs. full hybrid and BM25-only ($n = 199$). Both rows re-measured 2026-08-01 under the current retriever (both the exact-match coverage fix, Section 4.5, and the BM25 k_1/b retuning postdate the previous measurement of each row – see discussion below and Section 4.3).

Reranker	Comparison	Mean diff	p
On (fine-tuned)	vs. BM25-only, METEOR	−0.023	0.035
Off	vs. Full Hybrid, ROUGE-L	−0.017	0.205
Off	vs. BM25-only, METEOR	−0.025	0.041 (0.306)

(Section 4.3-era discussion below), and we re-ran it under the current retriever rather than assume it was unaffected. The adaptive pipeline (reranker off) is not significantly different from full hybrid on any of the four metrics ($p \geq 0.205$) or from vector-only on any of the four metrics ($p \geq 0.081$ across both tests) – a genuine flattening at the generation-quality level of what remains a clear retrieval-quality gap for vector-only specifically (Section 4.2), consistent with this paper’s broader pattern of retrieval-level and generation-level results not always moving together. Against BM25-only, the picture has changed from the previous revision: point estimates remain directionally behind on all four metrics, but none now reach clean significance under both tests. METEOR comes closest, marginally significant under the paired t -test alone ($p = 0.041$) while the Wilcoxon test does not agree ($p = 0.306$); consistent with how this paper treats other test-disagreeing results elsewhere (e.g. Section 3.12, and the previous revision’s own BERTScore result in this same comparison, now itself no longer significant under either test, $p = 0.097/0.808$), we treat this as inconclusive rather than confirmed. BLEU and ROUGE-L are directionally behind but not significant. In short, under the current retriever the adaptive pipeline is in full statistical parity with every single-method baseline it is compared against, not behind BM25-only on any metric at a confirmed significance level – a materially different picture from the previous revision’s confirmed METEOR loss, and we report the correction plainly rather than let the earlier, less favorable number stand uncorrected. It significantly outperforms no-retrieval on all four metrics ($p < 0.001$); this comparison excludes the 1 query (0.5%) on which the pipeline abstained – down sharply from 10.5% in an earlier measurement round, since the exact-match coverage fix (Section 4.5) and the resulting abstention recalibration (Section 3.7) together made the confidence gate fire correctly far less often on queries that are, in fact, answerable. An abstention is excluded rather than scored as a near-zero-BLEU miss, since that would conflate “declined to guess” with “answered wrongly” – two outcomes this paper’s abstention mechanism is specifically designed to keep distinct. We return to what this means for RQ1 in Section 5: on these metrics the pipeline does not clearly *beat* any single-method baseline either – parity is not automatically an advantage – and a deconfounding diagnostic (Section 4.7.2) confirmed the parity with BM25-only was not, on its own, evidence adaptive routing was a sound design choice: the router’s original fusion-method choice for entity-heavy queries (Reciprocal Rank Fusion) was directly shown to be significantly negative in isolation once a separate exact-match mechanism’s ranking ceiling was removed as a deconfounding step. That finding was acted on, not just reported: the fusion choice was replaced, verified to change no real-world behavior, and a genuine ranking-quality improvement was measured as a result (Section 4.7.2). What the pipeline adds is capabilities (abstention, graph augmentation, bilingual support) the single-method baselines do not have at all, not a generation-quality win.

4.4 Reranker Ablation

We also fine-tuned the reranker itself on this corpus’s own hard negatives – real retrieved near-misses from the live retriever, not random text – via a contrastive binary-cross-entropy objective, to test whether domain adaptation could flip its earlier net-negative result (Section 3.11) once training data actually came from the retriever it would be paired with. A first fine-tuning run happened to complete just before the exact-match coverage fix (Section 4.5) landed in the codebase, so its hard negatives reflected the OLD, pre-fix retriever’s candidate distribution rather than the one it would actually be deployed against – a genuine train/deploy mismatch we caught by comparing file timestamps rather than assuming the first result was clean, and which we flag here as a concrete pitfall for anyone fine-tuning a downstream component while the upstream retriever is still being iterated on. We re-ran the fine-tuning from scratch against the fixed retriever before drawing any conclusion.

Table 8 reports the fair, non-confounded comparison. The reranker-on row shown here is a second re-measurement, dated 2026-08-01: the version reported in an earlier revision of this paper (mean diffs of −0.040/−0.030 BLEU/METEOR vs. full hybrid, both $p < 0.03$, and −0.044/−0.030/−0.042 BLEU/ROUGE-L/METEOR vs. BM25-only, all $p \leq 0.016$) was measured before the exact-match coverage fix (Section 4.5) and the BM25 k_1/b retuning (Section 4.3) both landed, and we re-ran it rather than assume those upstream changes left it unaffected – the same discipline already applied once to this exact row when the embedding model changed. Under the current retriever, the reranker’s negative effect on this corpus is markedly weaker: it is no longer significantly different from full hybrid on any of the four metrics ($p \geq 0.221$) or from vector-only on any of them, and the only remaining significant difference against BM25-only is on METEOR (−0.023, $p = 0.035$ paired t , $p = 0.043$ Wilcoxon); BLEU, ROUGE-L, and BERTScore against BM25-only are all

Table 9: Unambiguous-match score-ceiling ablation, entity-heavy queries only ($n = 100$), re-measured 2026-08-01 after a second, later exact-match coverage fix (Section 4.5, the compound-query full-section-match bug) postdating the original measurement of this table; see discussion below.

Condition	Top-1 acc.	Top-5 acc.
Ceiling on (live)	0.930	0.930
Ceiling off	0.710	0.890

directionally behind but no longer significant ($p \geq 0.221$). We report this candidly as a real change in the measured margin, not a reversal to a positive result: the reranker still shows no advantage anywhere, and its one remaining loss (METEOR vs. BM25-only) is on the same comparison that is now only a borderline, test-disagreeing result for the deployed reranker-off pipeline too (Table 8, “Off” row, and Section 4.3), so the practical recommendation is unchanged – reranker-off remains the deployed default, since there is no metric on which enabling the reranker helps and its own METEOR-vs-BM25-only margin ($p = 0.035/0.043$, agreeing under both tests) is a more clearly confirmed loss than the reranker-off pipeline’s own version of the same comparison ($p = 0.041/0.306$, disagreeing).

As a further, literature-motivated check, we tested whether the reranker’s regression is a candidate-set-size effect: Magomere et al. (2025) find that reranking a wide candidate pool can inject noise that disproportionately harms an already-strong first-stage ranking, and recommend reranking a smaller pool as a direct mitigation. We re-ran the fine-tuned reranker restricted to a pool of 5 candidates (equal to `final_k`, so reranking can only reorder the already-selected top-5 rather than pull in noisier candidates from further down the ranking). This row has now been re-measured twice since the original pre-retrain test (2026-07-28, superseded): once on 2026-07-31 under the retrained embedding model alone, and again on 2026-08-01 under that same embedding model after the exact-match coverage fix (Section 4.5) and the BM25 k_1/b retuning also landed – we re-ran it a second time rather than assume the first re-measurement was still current, the same discipline applied to the pool-of-10 row above. Under the current retriever, restricting the pool does not rescue the reranker into a positive result, but it also no longer shows any measurable loss at all: pool-of-5 reranking is statistically indistinguishable from full hybrid, BM25-only, and vector-only on all four metrics ($p \geq 0.073$ in every comparison), including the two comparisons (BERTScore and METEOR vs. BM25-only) that were still significant in the previous re-measurement. Both directional differences remaining are small and mixed in sign (e.g. BLEU is directionally *ahead* of both full hybrid and BM25-only by +0.008 and +0.005 respectively, neither significant). We report this as a genuine further shrinking of an already-weakening effect, not a positive finding: with a restricted pool, this corpus’s reranker regression is no longer distinguishable from noise, so the informative comparison remains the pool-of-10 result above, where a real, if now smaller, negative effect persists on METEOR vs. BM25-only. Practically, this does not change the recommendation – reranker-off remains the deployed default, since restricting the pool earns parity at best, never an advantage.

4.5 Unambiguous-Match Score Guarantee, Isolated

Section 3.3 describes an absolute score ceiling applied when exactly one candidate is an unambiguous exact metadata match, guaranteeing it rank-1 rather than merely favoring it via an additive bonus. This mechanism had been live in every earlier round of this project but had not previously been isolated-tested on its own. A first isolated test, on the 100 English entity-heavy test queries, found the ceiling fires (i.e. exactly one exact-match candidate exists) for only 40 of the 100 queries. Inspecting the other 60 directly surfaced two concrete, fixable gaps rather than an inherent ceiling on the mechanism: 20 queries were faculty lookups (office room, email, designation) that the exact-match logic never recognized at all, since it only ever matched course-code patterns, even though the underlying FacultyList table (223 records) was already ingested; and 40 queries were “prerequisites for X” / “coordinator for X” shaped, where a blanket per-code CourseDetails cap (added to stop ~10–30 tied section rows from crowding out the one Prerequisites/Coordinator row a query actually wanted) was applied regardless of which table the query was actually asking about, so 2–5 candidates tied instead of one.

We fixed both: a faculty index (initials matched as standalone tokens, full names via normalized substring match) extends exact-match to FacultyList, and a table-type keyword check (“prerequisite”/“chain” → the Prerequisites row only; “coordinator” → the Coordinator row only) replaces the blanket per-code cap when the query names one specific table. After this fix, the same 100 queries have exactly one exact-match candidate in *all* 100 cases (previously 40) – both gaps were real coverage bugs, not queries the mechanism was never meant to handle. We then re-ran the isolated ceiling-on-vs.-off test from scratch under the fixed exact-match logic.

This table has itself been re-measured once since the version above: `scripts/test_unambiguous_match.py` predated a later, separate exact-match fix (the compound-query bug where a full-section match wrongly excluded its base code from Prerequisites/Coordinator lookup) as well as the BM25 k_1/b retuning. Under the current retriever, both conditions improved – ceiling-on rose from 0.850 to 0.930 and ceiling-off rose from 0.510 to 0.710 – narrowing but not closing the gap the ceiling itself contributes (+0.34 → +0.22 Top-1). This is the expected direction: the later fix

Table 10: FacultyAvailability test set, before and after the schedule-day routing fix ($n = 50$: 30 room, 20 day-schedule).

Query type (fix)	BLEU	ROUGE-L	BERTSc.	METEOR
Room (unaffected)	0.624	0.900	0.986	0.937
Day-schedule (before)	0.164	0.402	0.888	0.402
Day-schedule (after)	0.123	0.508	0.911	0.494

corrected a case the ceiling-off condition’s underlying exact-match logic was still getting wrong, so fixing it raises the floor ‘ceiling off’ measures against, without changing what the ceiling itself does. The qualitative conclusion is unchanged – the score-ceiling guarantee still contributes a large, real Top-1 gain on top of the additive bonus alone.

With every entity-heavy query now producing exactly one exact match, the ceiling’s marginal contribution to top-1 accuracy grows from the +8-point figure of the pre-fix mechanism to +34 points (Table 9); top-5 accuracy moves only marginally (+3 points), consistent with the additive bonus alone already being enough for top-5 presence in most cases, with the ceiling mainly deciding whether the correct row specifically reaches rank 1. Correctness for this test is judged against every distinctive entity in the reference answer (course codes and, newly, email addresses for faculty queries via regex extraction), with a name/initial-identity fallback for the room/designation queries that contain neither.

4.6 Faculty Lookup Coverage

FacultyAvailability (814 records: per-faculty, per-day teaching-schedule rows) was, as noted in Section 3.1, ingested and retrievable but not the source of any query in the automated-metric test sets above. We built a dedicated 50-query test set (30 office-room lookups, 20 day-schedule lookups, seed=42, reference answers composed directly from the table’s own fields) to close this gap directly, and in doing so surfaced a second, previously-untested exact-match coverage bug of exactly the same class as Section 4.5’s: the faculty exact-match index (built only from FacultyList, which stores office room/email/designation but no schedule data) forced its own row to rank-1 for schedule queries too, crowding out the actually-relevant per-day FacultyAvailability row entirely – confirmed live on “What is Ruhan Habib’s schedule on Sunday?”, which the pipeline answered “I do not have enough information” despite the corpus genuinely containing that day’s schedule. We fixed this the same way as the Prerequisites/Coordinator case: a schedule-and-day keyword check routes faculty exact-match to the FacultyAvailability row for the named day when the query asks about a schedule, and to FacultyList otherwise.

Room queries are unaffected by this fix, as expected, and score comparably to the main English test set. Day-schedule queries improve on three of four metrics after the fix (ROUGE-L +0.106, BERTScore +0.023, METEOR +0.092); BLEU alone moves slightly negative (−0.041), which we attribute to BLEU’s strict n -gram matching being sensitive to surface rephrasing rather than to a real regression – the qualitative change is from a non-answer (“I do not have enough information”) to a correct, differently-phrased one (*On Sunday, Ruhan Habib is scheduled: 8:00 AM – CSE321-17 (LAB) ...* scored against the pipeline’s own *Ruhan Habib’s schedule on Sunday includes a lab for CSE321-17 from 8:00 AM ...*), which BERTScore and METEOR credit and BLEU does not. We report the BLEU dip rather than omit it, consistent with this paper’s standing policy of reporting every metric’s actual movement.

4.7 English Lambda Sensitivity




Figure 5: Automated metrics vs. λ under the base (non-adaptive) linear fusion mode, on a fixed 60-query subset. No λ value significantly exceeds $\lambda = 1$ (BM25-only) on any metric (paired t -test).

Under plain (non-adaptive) linear fusion, sweeping $\lambda \in \{0, 0.1, \dots, 1.0\}$ finds no intermediate blend that significantly exceeds BM25-only on any of the four *generation-quality* metrics, in either the entity-heavy or open-ended subset (Figure 5); low-lexical-weight blends significantly *underperform* BM25-only, and performance is essentially flat across $\lambda \in [0.4, 1.0]$. This is the finding that directly motivated abandoning a single global λ in favor of the adaptive router (Section 3.5) rather than continuing to search for a better fixed blend point.

4.7.1 Is There Room for a Learned, Continuous Fusion Weight?

The routing decision above is still a binary choice between two fixed λ values (0.9 or 0.5); a continuous, learned, per-query fusion weight – in the style of Dynamic Alpha Tuning (DAT) proposals in the recent hybrid-retrieval literature – is a natural next idea, but before implementing one we checked whether this corpus’s queries actually disagree enough about their own best λ for such a weight to have anything to learn, using retrieval-only *nDCG@5* (a finer-grained, continuous IR metric than the generation-quality metrics above) on the 100 open-ended test queries, sweeping $\lambda \in \{0, 0.05, \dots, 1.0\}$ (21 values) per query.

Two things came out of this. First, the ceiling: an *oracle* that picks each query’s own single best λ in hindsight reaches mean *nDCG@5* ≈ 0.995 , against ≈ 0.98 for the single best fixed λ (0.25) applied to every query – a ceiling gain of only about 0.01–0.02 *nDCG@5* points from perfect per-query weighting, and only 32–35% of queries show *any* sensitivity to λ at all (the rest score identically at every λ we tried, i.e. retrieval either always finds or always misses the right chunk for them regardless of fusion weight). A real learned weight, unlike an oracle, would have to predict each query’s best λ from features available before generation, and could not exceed this already-small ceiling – so we did not implement DAT-style learned weighting: the expected gain, bounded above by roughly 0.02 *nDCG@5* points on a metric where two-thirds of queries are insensitive to λ by construction, does not justify the added model, training data, and maintenance surface a learned weight would add, especially set against this project’s other findings that added retrieval-stage sophistication (adaptive routing itself, Section 4.7.2) has repeatedly failed to beat simpler fixed configurations here. This ceiling was measured twice, once under each embedding model (before and after Section 3.4.1’s deployment): 0.017 *nDCG@5* points under the original embedding model, 0.011 under the mined-hard-negatives one – the same small-ceiling conclusion both times, which is why the decision not to implement a learned weight is reported as robust rather than tied to one specific embedding space.

Second, and more useful, a side effect of running this sweep: the single best fixed λ on this open-ended subset is 0.25, not the deployed default of 0.5, under both embedding models. $\lambda = 0.25$ significantly exceeds BM25-only ($\lambda = 1$) on *nDCG@5* in both measurements (original embedding model: mean diff +0.043, 95% CI [0.016, 0.073], $p < 0.0005$; hard-negatives embedding model: mean diff +0.046, 95% CI [0.022, 0.075], $p < 0.0005$) – a result that looks like it contradicts the generation-metric sweep above, but is measuring a different thing (ranking quality via *nDCG* vs. downstream generation quality via BLEU/ROUGE/BERTScore/METEOR) on a different-sized query set (100 vs. 60). $\lambda = 0.25$ also exceeds the deployed $\lambda = 0.5$ default in both measurements (original: mean diff +0.016, $p = 0.008$; hard-negatives: mean diff +0.019, $p < 0.0005$); the second measurement’s p -value comfortably survives a Bonferroni correction for the 21 λ values swept ($\alpha/21 \approx 0.0024$), where the first measurement’s did not. Two independent measurements (across a retrained, redeployed embedding model, i.e. not the same computation repeated) landing on the same best λ with the same direction of effect is meaningfully stronger evidence than either alone, though both still evaluate against the same 100-query set used throughout this paper, so we continue to report $\lambda = 0.25$ as a promising, now doubly-replicated lead for a future, properly held-out re-tuning of the open-ended default rather than a change we make unilaterally in this revision.

4.7.2 Isolating the Router’s Own Contribution

The lambda-sweep result above motivates routing but does not itself test the router: it shows no fixed blend beats BM25-only, not that the router’s specific choice for entity-heavy queries (Section 3.5) beats the alternatives on the queries where it actually takes effect. Pooling all 200 queries to compare adaptive against the full-hybrid baseline, as our first IR-metrics pass did, understates this by construction: `retrieve_adaptive`’s open-ended branch calls the exact same `retrieve(query, 0.5, fusion="linear")` as the full-hybrid baseline, so the two configurations are byte-identical, not just similar, on every open-ended query – any real effect can only appear on the 100/200 entity-heavy queries where adaptive’s own $\lambda = 0.9$ weighting takes effect instead.

Restricting the paired bootstrap comparison to exactly that 100-query entity-heavy subset (Table 11) gives a cleaner, and more deflating, answer than we expected: on Recall@1/3/5 and MRR, adaptive routing’s output is *identical*, row for row, to both the BM25-only and full-hybrid baselines – 100/100 ties on every one of these four metrics against both alternatives. This is not evidence that fusion choice doesn’t matter in general (Table 5’s full-corpus comparison still favors hybrid over BM25-only overall); it is evidence that our exact-match bonus (`EXACT_MATCH_BONUS` plus the unambiguous-match escalation, Section 4.5) already dominates ranking outright whenever an entity-heavy query resolves to a course code, faculty initial, or email that the corpus indexes directly – which the great majority of our

Table 11: Adaptive routing, isolated to the 100 entity-heavy test queries where it actually diverges from a fixed- λ baseline (paired bootstrap, 2000 resamples). Re-measured a second time (2026-08-01) after switching the entity-heavy route from RRF to linear fusion at the same $\lambda = 0.9$, a fix motivated by the first re-measurement’s finding (see text); the RRF-era numbers are reported in text, not repeated here.

Comparison	Metric	Adaptive	Other	p
vs. BM25-only ($\lambda = 1$)	Recall@1/3/5, MRR	0.930	0.930 (identical, 100/100 ties)	–
vs. BM25-only ($\lambda = 1$)	nDCG@5 / @10	0.883 / 0.909	0.892 / 0.915	<0.001 / <0.001
vs. Full Hybrid ($\lambda = 0.5$)	Recall@1/3/5, MRR	0.930	0.930 (identical, 100/100 ties)	–
vs. Full Hybrid ($\lambda = 0.5$)	nDCG@5 / @10	0.883 / 0.909	0.884 / 0.909	0.947 / 0.904
vs. Vector-only ($\lambda = 0$)	all 6 metrics	0.88–0.93	0.76–0.88	< 0.001 (all)

entity-heavy test queries do. Once that short-circuit fires, the top-ranked document is fixed regardless of the remaining candidates’ fusion: there is nothing left for the fusion method to decide.

This table has been re-measured twice since the version reported in an earlier revision. The first re-measurement (2026-08-01) reflected the compound-query exact-match coverage fix (Section 4.5): the tied Recall/MRR value moved from 0.850 to 0.930 (more entity-heavy queries now resolve to a genuine unambiguous match), and the nDCG picture, sensitive to the ordering of the otherwise-irrelevant remaining candidates, got measurably worse under the router’s own choice at the time (Reciprocal Rank Fusion, $\lambda = 0.9$): vs. BM25-only, nDCG@5 ($p = 0.030$) and nDCG@10 ($p < 0.001$) were both significantly behind, and vs. full hybrid both were non-significant but close to the threshold ($p = 0.060$, $p = 0.099$) – a real, if narrow, piece of evidence consistent with this paper’s deconfounding finding (Section 4.7.1) that adaptive routing’s own fusion choice, once anything is left for it to decide, tends negative rather than neutral.

Rather than leave that finding as a disclosed weakness, we acted on it: since RRF’s contribution is provably negative once isolated and provably *zero* in every real (non-isolated) test on this corpus – the exact-match ceiling fires on all 100/100 entity-heavy test queries regardless of fusion method – there is no case on this corpus where RRF outperforms simply using linear fusion at the same $\lambda = 0.9$ instead. We verified this directly before changing anything (0 discordant Recall@1 pairs between the two fusion choices under the live retriever, not assumed from the isolation test alone) and switched `retrieve_adaptive`’s entity-heavy branch to linear fusion. Table 11 reports the second re-measurement, under this fix: Recall/MRR ties are exactly preserved (0 regression, as verified), and the nDCG gap genuinely improved – vs. BM25-only it narrowed by roughly half ($-0.0217 \rightarrow -0.0087$ at nDCG@5) though it remains significant ($p < 0.001$ at both cutoffs); vs. full hybrid it collapsed to essentially zero (-0.0006 at nDCG@5, $p = 0.947$, no longer even borderline). RRF fusion (`fusion="rrf"`) remains implemented in the codebase for future use; this specific call site no longer selects it. We report this as a genuine, verified fix – not merely a disclosed limitation – while noting plainly that a small residual nDCG disadvantage vs. BM25-only remains and is not yet fully explained.

We re-ran this exact isolation after deploying the mined-hard-negatives embedding model (Section 3.4.1) and rebuilding the vector index, specifically to check whether the finding was an artifact of the previous embedding model rather than of the exact-match mechanism itself: it replicated exactly, 100/100 ties on Recall@1/3/5 and MRR against both fixed baselines, with the same small, non-significant nDCG gap. This is exactly what the underlying explanation predicts – the exact-match short-circuit doesn’t involve the embedding model at all – and finding the same result under an independently retrained embedding model is stronger evidence for the explanation than a single measurement could have been.

We report this as a genuine, if unflattering, isolated-ablation result rather than reframe it: on this corpus and test set, the adaptive router’s specific fusion-mode choice for entity-heavy queries contributes no measurable benefit over simply always using the full-hybrid default, because the exact-match mechanism – a much simpler, already-existing component – already does the actual work of getting these queries right. The router remains useful only insofar as its open-ended branch reproduces the already-validated full-hybrid default and its entity-heavy branch does no measurable harm; its added complexity (a second fusion mode, a second λ , RRF’s own untuned k constant, Section 4.7.3) is not currently justified by this evidence, and the honest conclusion is that a single fixed full-hybrid configuration, paired with the exact-match mechanism, would score identically on every metric we can measure. This does not mean routing is worthless in principle – only that its specific realization here has not been shown to add value beyond what a simpler system already provides, and we say so plainly rather than let the earlier pooled comparison’s smaller, insignificant-but-nonzero-looking numbers stand uncorrected.

The 100/100-tie result above is consistent with two different explanations: adaptive routing’s own fusion choice could be genuinely neutral (as good as the fixed baseline, just redundant given the exact-match mechanism), or it could be worse, with the exact-match mechanism’s rank-0 ceiling (`UNAMBIGUOUS_MATCH_SCORE`) simply masking a real deficit. We tested this directly rather than leave it ambiguous: temporarily patching `UNAMBIGUOUS_MATCH_SCORE` to 0 for a standalone diagnostic run only (`EXACT_MATCH_BONUS` left active; deployed behavior unaffected, the patch is reverted before the script exits), then re-measuring adaptive routing against the fixed full-hybrid baseline on the

same 100 entity-heavy queries. With the ceiling no longer available to rescue it, adaptive routing’s RRF-at- $\lambda=0.9$ choice (as originally deployed) was significantly *worse*, not neutral: Recall@1 0.71 vs. 0.93 (McNemar’s exact test, 22 discordant queries, all 22 favoring the fixed baseline, $p < 0.0001$), Recall@3 0.74 vs. 0.93 (19 discordant, all favoring the fixed baseline, $p < 0.0001$), Recall@5 0.88 vs. 0.93 (5 discordant, all favoring the fixed baseline, $p = 0.063$). Every single discordant query across all three cutoffs favored the fixed baseline; none favored adaptive routing. This resolved the ambiguity: the deployed system’s tie was not neutral redundancy, it was a worse fusion choice being propped up to parity by a separate, independently-validated mechanism. This finding is what motivated the fix reported in Section 4.7.2: rather than leave a component with a demonstrated negative effect and no offsetting benefit in place, we replaced RRF with linear fusion at the same λ for the entity-heavy route, verified the replacement changes no real (non-isolated) behavior, and measured a real improvement in the residual nDCG gap as a direct result. This diagnostic remains valuable after the fix precisely because it explains *why* that fix was safe to make: entity-collision handling (Section 4.5) is the real, positive mechanism doing the work; the specific choice of fusion method riding on top of it was not, which is exactly why removing that choice’s weakest option cost nothing.

4.7.3 RRF k Sweep: Confirming, Not Assuming, the Same Explanation

This sweep was run while RRF was still the deployed entity-heavy fusion method, and is retained as supporting evidence for the fix in Section 4.7.2, not as an active tuning question for the current system. RRF’s own rank-damping constant ($k = 60$ in `_score_rrf`, following the original RRF paper’s conventional default (Cormack et al., 2009)) had never been tuned for this corpus. If the exact-match short-circuit explanation above is correct, sweeping k should show *no* effect on Recall/MRR for entity-heavy queries specifically, since the forced rank-0 assignment plus the unambiguous-match bonus fixes the top document before k is ever consulted; if that explanation is wrong or incomplete, some sensitivity to k should remain. We swept $k \in \{10, 20, 40, 60, 100, 200\}$ (a $20\times$ range) on the same 100 entity-heavy queries, and first checked how many of them actually have an unambiguous (single) exact match, since only the remainder could show any k -sensitivity at all: all 100/100 do. Recall@1, Recall@5, and MRR were exactly 0.850 at every single value of k across this entire range – not approximately flat, identically flat, to four decimal places, confirming the mechanism rather than merely being consistent with it. nDCG@5/@10 showed small, non-monotonic fluctuation (nDCG@5 ranged 0.792–0.799 across the sweep, from re-ordering of already-irrelevant lower-ranked candidates below the fixed top slot). This result – that k has no real influence because RRF is never the deciding factor once the exact-match ceiling fires – is itself part of the case that replacing RRF with linear fusion (Section 4.7.2) could not lose anything the deployed system actually relied on.

The underlying reason RRF’s contribution could never be observed on this test set is a dataset limitation, not a settled fact about RRF in general: every entity-heavy query in our 200-query test set happens to resolve to an unambiguous exact match. We have no queries in the test set where entity resolution is ambiguous (multiple candidate matches) or fails outright, which is exactly the regime where a rank-fusion method’s actual contribution – as opposed to the exact-match bonus riding on top of it – could be measured, and where RRF might have shown a genuine advantage over linear fusion had one existed. Constructing such a subset (e.g. deliberately ambiguous faculty-name queries, or course codes shared across cross-listed sections) is future work; RRF fusion remains implemented and available in the codebase (`fusion="rrf"`) for that subset to be tested against once it exists, even though it is no longer the entity-heavy route’s default.

4.8 Faithfulness (Automated LLM-Judge)

Automated similarity metrics cannot, by construction, distinguish a faithful paraphrase from a confidently stated factual error. Rather than a manual annotation study, which for this revision we do not report (Section 5 discloses why), we adopt RAGAS’s faithfulness formulation (Es et al., 2023): the same locally hosted LLM used for generation is separately prompted to decompose a generated answer into its individual factual claims and score, from 0 to 1, the fraction actually supported by the retrieved context it was given – not whether each claim is true in general, only whether the provided context backs it up. We score a stratified sample of English outputs: 40 per baseline configuration and 50 for the adaptive pipeline (reranker off), on the identical query_ids re-measured 2026-08-01 under the current retriever (the original sample predated the exact-match fix, BM25 retuning, and the Banglish-expansion embedding retrain, the same staleness class as every other table in this revision).

Both retrieval-based baselines still score highest (0.916–0.923), the adaptive pipeline still scores lower but well above no-retrieval (Table 12). We report this pattern as-is rather than reconcile it into a stronger claim: it is not a paired, significance-tested comparison at face value, since the samples are drawn independently per configuration at different sizes – Section 5’s Limitations report the actual paired, matched-subset test separately. One descriptive change from the previous measurement is real and worth naming: Vector-only’s score rose noticeably (0.761→0.858), consistent with the significant embedding-quality improvement from the Banglish-expansion retrain measured directly elsewhere (Section 3.4.1) – this is not noise, it is the same underlying cause showing up in a second, independent metric.

Table 12: Automated faithfulness score (LLM-judge, 0–1), English sample, re-measured 2026-08-01 on the identical query_ids under the current retriever.

Configuration	Mean faithfulness	n
Full Hybrid	0.923	40
BM25-only	0.916	40
Vector-only	0.858	40
Adaptive pipeline	0.810	50
No-retrieval	0.595	40

Table 13: Independent NLI-based faithfulness cross-check (sentence-level, SummaC-ZS-style), vs. the LLM-judge score, re-measured 2026-08-01 on the same regenerated sample as Table 12.

Configuration	LLM-judge	NLI (independent)
Full Hybrid	0.923	0.562
BM25-only	0.916	0.516
Adaptive pipeline	0.810	0.519
Vector-only	0.858	0.494

We investigated, rather than left unexplained, what distinguishes the adaptive pipeline’s lower-scoring rows. Four candidate explanations were checked directly and each ruled out, consistent with the previous measurement: mean generated-answer length (21.3 words for the adaptive pipeline vs. 20.6–22.8 for the retrieval baselines – not meaningfully longer, ruling out a simple verbosity effect; no-retrieval’s 53.1 words is a separate, expected effect of the model having no grounding to constrain it); route (entity-heavy 0.84 vs. open-ended 0.79 within the adaptive sample – too small a gap to explain the difference from baselines); graph augmentation (0.825, $n = 4$, vs. 0.809, $n = 46$ – too small a subsample to be informative); and exact-match status (0.84 vs. 0.79 – similarly small). Qualitative inspection of the lowest-scoring rows continues to show the same recurring, specific pattern identified in the original measurement: the judge’s claim-by-claim decomposition sometimes does not credit a natural-language restatement of information that is genuinely present in the assembled context but in a different form – e.g. a paraphrased prerequisite-graph chain (Section 3.6) that the judge’s reason text flags as “not explicitly stated” despite the graph block being part of the recorded context. This is a real, disclosable, partially-explanatory finding – a likely interaction between the richer, more structurally-assembled context the adaptive pipeline provides (graph blocks, additional pool-matched chunks) and a strict per-claim automated judge – but we do not consider the cause fully isolated, and list further investigation (ideally with a second, independently-worded judge or human annotation) in Section 5.

4.8.1 An Architecturally Independent Cross-Check

The faithfulness score above has a structural limitation beyond sample size: the judge is the same model that performs generation, so it is not an independent rater in any meaningful sense. We built a second, architecturally unrelated faithfulness signal using an off-the-shelf natural-language-inference (NLI) model never trained on this corpus, motivated directly by the HALT-RAG line of work (Goswami and Kurra, 2025): does the retrieved context entail the generated answer? A naive first attempt – one NLI call per row, premise = the full retrieved context, hypothesis = the full generated answer – produced near-meaningless scores (0.20–0.25 across all configurations) with essentially zero correlation to the LLM-judge score (Pearson $r = 0.034$, $n = 170$), which we take as a red flag rather than a finding: standard NLI cross-encoders are trained on short, single-sentence pairs and are known to degrade sharply on long, multi-claim premises and hypotheses. We corrected this using the SummaC-ZS approach (Laban et al., 2022): split both context and answer into sentences, and for each answer sentence take the maximum entailment probability over all context sentences, averaging across answer sentences per row.

The two methods’ absolute scales are not comparable – NLI entailment is a strictly narrower criterion than the judge’s rubric, so uniformly lower NLI scores are expected and not evidence of worse grounding – and row-level agreement between them is, if anything, weaker under this re-measurement than before (Pearson $r = -0.063$, binary agreement at a 0.5 threshold = 49.4%, indistinguishable from chance, $n = 170$). This re-measurement changes the previous revision’s central NLI finding, and we report the correction directly rather than let a superseded interpretation stand: the independent NLI check no longer clearly fails to reproduce a faithfulness gap for the adaptive pipeline. Adaptive (0.519) remains essentially tied with BM25-only (0.516, matching the previous measurement’s pattern), but is now descriptively *behind* full hybrid (0.562, a -0.043 gap) rather than slightly ahead of it as the previous measurement showed – the opposite of the pattern the claim-decomposition-bias hypothesis predicted. The gap under NLI (-0.043) is still substantially smaller than the LLM judge’s own gap against full hybrid (-0.113 in raw means; see the matched-subset test below for the paired figure), so this is not a case of the independent check fully agreeing with the judge

Table 14: Banglish retrieval ablation and adaptive pipeline, re-measured 2026-08-01 under the deployed embedding model (`finetuned_minilm_hard_negatives_structured_banglish_expanded`) and the current retriever ($n = 100$). The previous measurement (2026-07-28) predated this embedding model’s own training run (2026-07-31) as well as the exact-match and BM25 fixes, and was found stale on re-verification – see discussion below.

Variant	BLEU	ROUGE-L	BERTSc.	METEOR
Full Hybrid	0.822	0.900	0.979	0.900
Adaptive pipeline	0.811	0.899	0.980	0.902
BM25-only	0.825	0.888	0.978	0.891
Vector-only	0.830	0.907	0.981	0.900
No-retrieval	0.015	0.134	0.832	0.184

either – but with row-level agreement between the two methods now indistinguishable from chance, we no longer treat this comparison as corroborating evidence for the claim-decomposition-bias hypothesis in either direction. The honest reading is that this specific cross-check is too noisy, at this sample size and with this NLI model, to confirm or rule out the hypothesis, not that the hypothesis itself is confirmed or refuted by it.

4.9 Banglish Evaluation

Table 14 has itself been re-measured since the previous revision, and the reason turned out to be different from every other staleness fix in this paper: unlike the exact-match and BM25 fixes, which only affect configurations with $\lambda > 0$, the previous Banglish measurement (2026-07-28) predated the Banglish-expanded hard-negative retrain (2026-07-31) that is the embedding model actually deployed today (Section 3.4.1) – a staleness source specific to this table, since Vector-only ($\lambda = 0$) is the one configuration purely dependent on embedding quality and cannot be affected by either the exact-match or BM25 fixes at all. Re-scoring against the previous measurement confirms this directly: Vector-only improved significantly on all four metrics (BLEU +0.057, ROUGE-L +0.045, BERTScore +0.008, METEOR +0.039, all $p < 0.05$, paired t -test, $n = 100$), while BM25-only and Full Hybrid show no significant change (BM25-only BLEU $p = 0.058$, borderline but not significant; all other comparisons $p \geq 0.13$), consistent with those two configurations depending on the lexical stream and the exact-match mechanism far more than on embedding quality specifically. This is a genuine, previously-undiscovered staleness gap we are glad to close, and a real positive result for the Banglish-expanded retrain: it was simply never re-measured against the full ablation table after that retrain landed.

With the current embedding model in place, matched paired testing ($n = 98/100$, two queries abstained on, up from one in the previous measurement) finds the adaptive pipeline statistically indistinguishable from full hybrid, BM25-only, and vector-only on all four metrics ($p \geq 0.31$ throughout under both tests) – notably cleaner than the previous measurement’s one borderline exception (BERTScore vs. BM25-only, then $p = 0.046/0.116$, now $p = 0.955/0.548$, no longer close to significance under either test) – and significantly ahead of no-retrieval ($p < 0.001$). This remains a cleaner tie than the English result: on this test set, retrieval method essentially does not matter for Banglish queries as long as some form of retrieval is used at all, which we discuss further in Section 5 together with a structural reason to treat this test set’s ability to discriminate between retrieval methods with some caution.

4.9.1 Query Translation (Null Result)

Comparing the adaptive pipeline with cross-lingual query translation (Section 3.12) enabled vs. disabled, on the same 100 Banglish queries: BLEU 0.810 vs. 0.819 ($p = 0.68$), ROUGE-L 0.889 vs. 0.899 ($p = 0.46$), BERTScore 0.978 vs. 0.979 ($p = 0.60$), METEOR 0.890 vs. 0.892 ($p = 0.91$). None of the four differences is significant, and the point estimates slightly favor *not* translating. Query translation is implemented and available but disabled by default.

4.9.2 BM25 Code-Mixed Normalization (Null Result)

Comparing BM25-only and full-hybrid retrieval with the Banglish spelling-variant normalization layer (Section 3.12) enabled vs. disabled: for BM25-only, all four metrics move in the negative direction with normalization on but none reach significance ($p = 0.13$ – 0.54); for full hybrid, differences are smaller and also not significant ($p = 0.42$ – 0.91). Normalization is implemented, applied consistently at index-build and query time, and left in the tokenizer, but we do not claim it helps on this test set.

4.9.3 Why the Two Banglish-Specific Ablations Show Null Effects

Both null results above share a structural explanation. This Banglish test set is sampled from BanglishQA’s own stored questions, whose spelling and phrasing already match the corpus’s own indexed content by construction (the query *is* a held-out copy of a question the corpus already contains, just from the test rather than train partition). Query translation

Table 15: Cross-lingual stress test: top-5 retrieval accuracy, query translation on vs. off ($n = 13$: the original 9 test-split facts plus 4 further val-split facts, both held-out splits, EnglishQA-only content, Banglish-phrased queries). Re-measured 2026-08-01 under the current retriever.

Condition	Top-5 accuracy
Plain retrieval (no translation)	0.615 (8/13)
With query translation	0.846 (11/13)

and lexical normalization are both designed to help when the query’s language or spelling *diverges* from how the relevant content happens to be indexed – a scenario this test set structurally does not probe, since there is no divergence to bridge. Section 4.9.4 tests this explanation directly rather than leaving it as a hypothesis.

4.9.4 Cross-Lingual Stress Test: Query Translation, Validated

To test the explanation above directly, we built a small stress test targeting exactly the scenario query translation and normalization are designed for and the main Banglish test set structurally cannot probe: Banglish queries whose relevant content exists *only* in EnglishQA, not BanglishQA. We identified EnglishQA test-split rows whose answer text has no match anywhere in BanglishQA’s own answer field (21 of 244 EnglishQA test rows qualify; deduplicating by unique answer text, since several are paraphrase duplicates of the same fact, yields 9 rows), then rephrased each question into natural Banglish using the same local LLM used for generation – disclosed here as the construction method, in the spirit of standard code-switched-NLP dataset construction practice (LLM/MT-assisted rephrasing), not presented as organically collected data. We then extended this with the same procedure applied to the val split (never trained on, the same held-out status as test for this purpose): 8 more EnglishQA-only rows, deduplicating to 4 further unique facts. We do *not* draw on the train split for this expansion – those rows were seen directly during embedding fine-tuning (Section 3.4), so including them would leak training data into a held-out stress test. The combined $n = 13$ set is what Table 15 reports.

Unlike the null result on the main Banglish test set, query translation still improves top-5 accuracy here (Table 15), directionally confirming the structural explanation of Section 4.9: translation was never an ineffective technique, it was being tested on a set with no genuine cross-lingual gap for it to bridge. Both this table and its predecessor have been re-measured since the previous revision, and the change is informative rather than incidental: the plain-retrieval baseline rose substantially from an original 0.333 (at the original $n = 9$) to 0.556 and now 0.615 at $n = 13$, while the translated condition has stayed close to its original level throughout ($7/9 \rightarrow 8/13$, i.e. still ≈ 0.85) – the BM25 retuning, exact-match fix, and embedding retrain (all Section 4.3-era changes) improved plain cross-lingual retrieval enough to partly close the gap this stress test was built to detect, narrowing but not eliminating translation’s advantage. The val-split addition is a real replication check, not just a bigger number: the gap held up under 4 entirely new facts the earlier $n = 9$ measurement never saw ($+0.231$ at $n = 13$ vs. $+0.222$ at $n = 9$), which is modest evidence the effect is not an artifact of the original small sample. It still does not reach significance (exact binomial test on the 5 discordant pairs, 4 favoring translation and 1 favoring plain retrieval, $p = 0.375$) – worth naming plainly why: at only 5 discordant pairs, even a unanimous 5/5 split would not clear $p < 0.05$ (binomial $p = 0.0625$ in that case), so this specific test is structurally underpowered to reach significance at any outcome until the discordant-pair count itself grows substantially, which requires more genuinely novel EnglishQA-only facts than currently exist in the corpus, not a different statistical test. We report this as a suggestive, mechanistically-explained, now twice-replicated result rather than a significance-tested one – a larger version of this specific stress test, as the corpus grows, remains future work (Section 5) – and the direction continues to support the earlier correction of query translation’s framing from simply unhelpful, though the margin is smaller than first measured.

4.9.5 Cross-Lingual Sufficient-Context Gating: A Mechanistic Diagnostic

The retrieval-accuracy result above shows THAT translation helps on this stress set; we additionally tested WHY, at the level of individual queries, by applying Joren et al. (2024)’s sufficient-context classification – could a diligent reader answer using only this context? – twice per query: once against context retrieved via the original Banglish phrasing, and once against context retrieved via its English translation. To our knowledge this bilingual application of sufficient-context classification, used to diagnose translation-sensitive retrieval failures specifically, has not been reported in either the sufficient-context or the code-mixed-retrieval literature we surveyed. Four outcomes are possible: both judged sufficient (translation adds nothing new); both insufficient (a genuine retrieval gap, not a language-routing issue); original-language context judged insufficient but translated-language context sufficient (*translation-sensitive sufficiency* – direct, per-query evidence of translation earning its keep); or the reverse (translation making things worse).

Two of the nine queries show the translation-sensitive pattern directly, and none show the reverse (Table 16) – concrete, per-query mechanistic confirmation of the aggregate retrieval-accuracy finding above, rather than a second independent claim. With $n = 9$ this is illustrative rather than statistically powered, but it demonstrates a diagnostic this

Table 16: Cross-lingual sufficient-context outcomes ($n = 9$, same stress-test queries as Table 15).

Outcome (original / translated)	Count
Sufficient / Sufficient	6/9
Insufficient / Insufficient	1/9
Insufficient / Sufficient (translation-sensitive)	2/9
Sufficient / Insufficient	0/9

Table 17: This paper’s ablations regrouped by RAG pipeline stage (Huang and Huang, 2026). Δ is relative to the strongest single-method baseline (BM25-only or full hybrid, whichever applies).

Stage	Component tested	
Pre-retrieval	Query translation (Banglish)	Null on main s
Pre-retrieval	Entity normalization	Regex fix confi
Retrieval	Embedding fine-tuning, mined hard negatives	Significant gain over in-batch-only fine
Retrieval	Fusion choice (BM25/dense/hybrid)	Hybrid > BM
Retrieval	Adaptive routing vs. fixed λ	Isolated to entity-heavy queries (Sec. 4.7.2): original RRF choice f
Retrieval	Learned/continuous fusion weight (DAT-style)	Not implemented: oracle ceiling only +0.01–0.0
Retrieval	Unambiguous-match ceiling	+22pp top-1 accu
Post-retrieval	Cross-encoder reranking	Negative but weaker on re-measu
Post-retrieval	Context ordering (zig-zag)	Implemen
Architecture	Prerequisite-graph augmentation	Isolated, expanded to $n = 31$ real DB-verified
Architecture	Confidence-gated abstention	Abstention rate 10.5% $\rightarrow$ 0.5% after retriever/calibra
Generation	Faithfulness (LLM-judge)	Adaptive lower descriptively; matched-subset test sig

project’s own tooling can already produce cheaply (two additional LLM calls per query, no fine-tuning) to characterize exactly when cross-lingual translation matters for a given query, rather than only reporting that it helps in aggregate.

4.9.6 Expanding the Stress Test: A More Cautious Picture at Larger n

Since the stress test’s bottleneck is genuinely EnglishQA-only *facts* (9 unique ones in the test split), not phrasing diversity, we expanded it by generating three additional Banglish phrasings of each of the same 9 facts via the same local LLM, using temperature 0.7 rather than this project’s otherwise-universal temperature 0.0 specifically to obtain distinct rephrasings (greedy decoding would return an identical rephrasing on every repeated call) – yielding 27 queries over the same 9 underlying facts. Following Doğruöz et al. (2023)’s caution that LLM-generated code-switched text can look more “textbook” than naturally-occurring code-switching and may not represent how real bilingual speakers actually mix languages, we disclose this directly rather than treat the expanded set as equivalent to the original, more conservative one: manual inspection of the generated variants found visible quality problems in a minority of them, including one case where the translation step’s own output leaked meta-commentary (“*However, this is not a question that needs to be translated...*”) instead of a clean translation.

On this expanded, noisier set, re-measured 2026-08-01 alongside Table 15 under the current retriever, the same pattern as the $n = 9$ result appears at larger scale and even more attenuated: plain top-5 accuracy 59.3% (16/27, up from 40.7% previously, the same retriever-improvement effect as the smaller set) vs. translated 63.0% (17/27, essentially unchanged from 55.6% previously); of the 5 queries where the two conditions disagree, 3 favor translation and 2 favor plain retrieval (exact binomial test on discordant pairs, $p = 1.0$, compared to $p = 0.219$ previously). We report this honestly as a further-weakened picture rather than omit it: the direction still nominally favors translation, but with the plain-retrieval baseline now much stronger on its own, this expanded set no longer offers even suggestive statistical support, only the original, more conservative $n = 9$ set (Table 15) still does. Statistical confirmation requires either a larger genuinely-novel-fact set (not more phrasings of the same 9 facts) or human-validated generated variants, neither of which this revision had time to produce – and, as of this re-measurement, a corpus and retriever where the underlying task is not already partially solved by improvements orthogonal to translation itself.

4.10 Results Organized by Pipeline Stage

The results above are reported per-experiment, in the order they were run; Huang and Huang (2026)’s organization of RAG systems into pre-retrieval, retrieval, post-retrieval, and generation stages offers a complementary view of the same evidence, showing which stage each component’s measured effect actually belongs to rather than a flat list of ablations. Table 17 regroups this paper’s own results by stage; every number in it is reported elsewhere in this section and is repeated here only for this organizational view, not re-measured.

Presented this way, a pattern is visible that a flat list of ablations does not make as obvious: this system’s clearly positive, measured effects are the exact-match mechanism (Table 9) and, now confirmed at $n = 31$, prerequisite-graph augmentation (Table 2), while every post-retrieval intervention tested (reranking, in two forms) was negative. Prerequisite-graph augmentation, the one architecture-level component isolated so far, reversed its own formatting-artifact regression, then cleared this paper’s full significance bar once measured on a large enough real sample of the chains the corpus actually supports (Sec. 3.6). Adaptive routing, once isolated to the queries where it actually diverges from a fixed baseline (Sec. 4.7.2), turns out to add no measurable value on top of the exact-match mechanism and full-hybrid fusion alone – a concrete instance of a broader lesson this table makes visible: several components motivated by reasonable-sounding evidence (a lambda sweep, a reranker’s typical benefit in the literature) do not survive isolated, fairly-paired testing on this specific corpus, while the simplest mechanism (direct entity matching) is the one component nothing else has managed to beat.

4.11 Pipeline Efficiency

We evaluated the ingestion pipeline under three operating conditions. A cold-start ingestion of the full corpus required 368.62 seconds. Re-ingesting after updating the underlying source files required 106.82 seconds. Re-ingesting unchanged data required only 4.37 seconds. Figure 6 summarizes this comparison; these timings predate the corpus and pipeline extensions in this revision and are reported as an approximate, not re-measured, reference point.




Figure 6: Comparison of data ingestion time across cold-start, incremental-update, and no-change scenarios.

5 DISCUSSION

On RQ1, under the current retriever the adaptive pipeline does not clearly *beat* BM25-only on any of the four automated similarity metrics, but nor is it confirmed behind it on any of them either – it is statistically tied on BLEU and ROUGE-L, and, as of the current re-measurement, only marginally and inconclusively behind on METEOR (significant under a paired t -test alone, $p = 0.041$, not under a matched Wilcoxon test, $p = 0.306$) and not significant at all on BERTScore (Section 4.3) – and remains statistically tied with full hybrid and with vector-only on all four. This supersedes two earlier findings in turn: an original pre-retrain measurement found a two-metric edge for the adaptive pipeline, which traced to a stale measurement never re-verified after the embedding model changed; the corrected measurement then found a confirmed METEOR loss, which itself went stale after the exact-match coverage fix and BM25 retuning and has now been re-verified a second time. We report each correction plainly rather than let a superseded number stand. What the pipeline adds is not a generation-quality win but capabilities none of the single-method baselines have at all: abstaining rather than guessing on the rare open-ended English query it cannot ground with adequate confidence (down to 0.5% of queries after the exact-match and abstention-recalibration fixes below, from 10.5% in an earlier revision); correctly answering multi-hop prerequisite-chain questions via graph traversal, a query shape no single retrieved chunk answers; and operating over a bilingual corpus with the same architecture and comparable retrieval-ablation results in both languages. On RQ2, direct ablation shows a mixed picture by component: fine-tuning the embeddings (Section 3.4), the unambiguous-match score guarantee (Section 4.5), and prerequisite-graph augmentation (Section 3.6, confirmed on all four metrics once measured on an expanded $n = 31$ sample of the corpus’s own real prerequisite chains, not just the 12 the main test set happens to exercise) each produced a measured, positive effect; adaptive routing was directly tested, not merely inferred from the λ -sweep finding that motivated it, and its original design (RRF fusion for entity-heavy queries) was found to be actively counterproductive on its own terms – isolated to the entity-heavy queries where it can differ from a fixed baseline, it tied the baseline exactly under the deployed exact-match mechanism (Section 4.7.2), but with that mechanism’s ranking ceiling removed as a deconfounding diagnostic, the RRF choice specifically was significantly worse ($p < 0.0001$ on Recall@1/3), showing the tie was a mechanism

outside adaptive routing rescuing it, not evidence the RRF choice itself was sound. We replaced it with linear fusion at the same weighting once this was established, confirmed the replacement preserves the tie exactly (0 regression) while narrowing the residual nDCG gap; and the cross-encoder reranker produced a measured, significant *negative* effect, which we confirmed is not simply an artifact of an undertrained reranker (fine-tuning it on the corpus’s own hard negatives from the actual deployed retriever did not flip this result) nor of candidate-set noise from a wide reranking pool (restricting the pool to `final_k=5` candidates did not flip it either, Section 4.4) – on this corpus, the reranker’s joint relevance judgments disagree with the fine-tuned embeddings’ ranking under both explanations the literature offers for when reranking helps or hurts. On RQ3, the system performs comparably across languages on the retrieval ablation, and the two components built specifically for the code-mixed setting (query translation, BM25 normalization) show null effects on the main Banglish test set for a structural reason we then tested directly: a dedicated cross-lingual stress test (Section 4.9.4), built from EnglishQA content with no BanglishQA counterpart, shows query translation improving top-5 retrieval accuracy (61.5% to 84.6% on re-measurement, $n = 13$ after a legitimate expansion with 4 further real, held-out-split facts) in the scenario it was actually designed for – confirming the null result was at least partly a test-set limitation, not solely evidence the technique is unhelpful, though this revision’s own retriever improvements narrowed the margin considerably (the plain-retrieval baseline rose from 33.3% to 61.5% on the same underlying set) and the larger, noisier expanded version of this stress test no longer reaches significance on the current retriever, so we hold this conclusion more cautiously than the previous revision did.

We describe this paper’s contribution deliberately as a validated engineering integration, not a novel retrieval algorithm: Reciprocal Rank Fusion (Cormack et al., 2009), contrastive embedding fine-tuning (Henderson et al., 2017), cross-encoder reranking, graph-based multi-hop reasoning, confidence-gated abstention with a sufficiency override (Joren et al., 2024), and confidence-ordered zig-zag context assembly (Liu et al., 2024; Jin et al., 2024) are each individually established techniques. What we report as the contribution is that we combined them into one system for an under-studied bilingual academic-advising setting, ablated each component individually rather than only end-to-end, and report the resulting evidence – including where components did not help – rather than assume the combination is an improvement because each part is individually well-motivated. Where we come closest to a genuine unifying idea rather than a combination is Section 3.9’s observation that abstention gating, context ordering, and (as future work) compression can all be driven by the same per-chunk confidence score computed once, and Section 4.9.5’s bilingual application of sufficient-context classification as a diagnostic for code-mixed retrieval specifically – both modestly scoped, and reported as such.

5.1 Limitations

Several limitations qualify these findings, some resolved during this revision and some still open. First, we still have not produced a manual hallucination annotation; the automated LLM-judge faithfulness score (Section 4.8) is a real, code-executed measurement but is a different instrument from human annotation against ground truth, uses a single automated judge with no second rater, and its lower point estimate for the adaptive pipeline relative to the baselines is only partially explained (Section 5). This is the one limitation on this list that cannot be resolved by more engineering effort alone – it requires an actual second human rater, which is outside the scope of what this revision can produce. Second, the entity-heavy abstention threshold (Section 3.7) was calibrated against only 18 naturally-occurring labeled examples; we grew this to 60 with database-verified synthetic examples (nonexistent vs. real course codes/faculty initials), tightening the 95% CI on accuracy from (0.65, 0.99) to (0.861, 0.990) – better, but the synthetic portion is template-generated rather than organic, which we disclose rather than blend in silently. Third, the open-ended abstention threshold’s recall improved from 0.457 to 0.870 after recalibrating against the exact-match-fixed retriever (Section 4.5); the sufficient-context override (Section 3.7) still only ever makes the system *less* likely to abstain, so it cannot itself compensate for missed detections, but the underlying gate is now substantially more sensitive than before. A deeper methodological gap in this same calibration, found and closed this revision: the threshold-selection procedure sweeps for the accuracy-maximizing threshold on the same set it then reports accuracy on, with no train/test split, so the reported figures above are in-sample fit quality, not a genuine generalization estimate. A held-out re-validation (stratified 70/30 train/test split, seed 42, threshold selected on train only and evaluated on a test set never used for selection) finds this gap is real and asymmetric: the entity-heavy route generalizes well (train accuracy 1.000, held-out test accuracy 0.944, $n_{\text{test}} = 18$, close to the in-sample figure above), but the open-ended route does not (train accuracy 0.673, held-out test accuracy 0.532, $n_{\text{test}} = 111$, barely above chance on this class-balanced set) – substantially weaker than the 0.710 in-sample accuracy this section otherwise reports. We disclose this as a genuine, previously-unmeasured weakness of the deployed open-ended abstention gate’s generalization, not merely its known small-sample-CI limitation; the sufficient-context override partially mitigates the practical impact (a missed abstain-worthy query can still be caught by structural evidence when present), but the confidence threshold itself is weaker out-of-sample than in-sample reporting had suggested. A single 70/30 split is itself sensitive to which points happen to land in the test fold, so we went further and actually re-calibrated with 5-fold cross-validation (threshold selected on 4 folds, evaluated on the held-out fifth, for all 5 folds; final deployed threshold and signal chosen by majority vote across folds, threshold set to the median of the winning folds’ thresholds): the open-ended route’s cross-validated mean accuracy is 0.612 (std 0.017

across folds, a tight, stable estimate) – notably higher than the single-split estimate above, showing that split was itself somewhat pessimistic, but still meaningfully below the 0.710 in-sample figure originally reported. The resulting threshold (0.9916, close to but lower than the previous 1.0283) is now deployed (`results/abstention.threshold.json`), replacing the in-sample-overfit selection with a cross-validated one, verified live to behave sanely on real queries before deployment. The entity-heavy route’s cross-validated selection surfaced a different, more subtle issue worth disclosing rather than acting on blindly: its `query_confidence` signal is sharply bimodal on the calibration set (every value falls in $[0, 0.011]$ or $[99.98, 100.0]$, with no examples in between, a direct consequence of the unambiguous-match ceiling mechanism, Section 4.5), so a threshold anywhere in this wide empty gap scores identically on the calibration data – cross-validation picked a threshold near the low end (0.0114) purely because of this degeneracy, not because it is a better choice. We did not deploy it: a threshold near the ceiling (the existing 99.9881) is the more conservative, safety-consistent choice for a genuinely novel intermediate-confidence query that this calibration set has never seen, even though the two thresholds are statistically indistinguishable on the data available. This is a concrete illustration of why cross-validated selection should not be deployed automatically without a sanity check against the underlying mechanism it is calibrating. Fourth, the cross-lingual stress test identified as a gap in an earlier pass – Banglish queries about content that exists only in EnglishQA – has now been built and evaluated (Section 4.9.4), and legitimately expanded once, from the original 9 test-split facts to 13 by adding 4 further val-split facts (never trained on, the same held-out status as test); the direction replicated under the addition. Its main limitation remains size: only 13 qualifying real facts exist across both held-out splits combined (the train split cannot be used without leaking training data into the stress test), so its result is reported as suggestive rather than significance-tested, and we show plainly why – at this few discordant pairs, even a unanimous outcome could not clear conventional significance. Fifth, the faithfulness sample sizes (40–50 per configuration) remain modest relative to the full 200-query English test set; we did apply paired bootstrap significance testing to them following Dror et al. (2018)’s guidance that graded, potentially non-normally-distributed measures like faithfulness scores should use a non-parametric paired test rather than an assumed-normal paired t -test, but this was only possible in full among the four baseline configurations (confirmed to share the same 40 query_ids); the adaptive-pipeline sample only overlaps the baselines’ query_ids on 17/40 rows, restricting that comparison to a smaller matched subset. Re-measured 2026-08-01 on the regenerated sample (Section 4.8), this matched-subset test now shows a more nuanced picture than the previous revision’s uniform null result: the adaptive pipeline’s faithfulness is significantly *higher* than vector-only ($+0.112$, $p < 0.001$) and no-retrieval ($+0.299$, $p = 0.003$) on this subset – a sanity check that now actually passes, unlike the previous measurement where even these comparisons were non-significant – while remaining statistically indistinguishable from BM25-only ($+0.055$, $p = 0.26$) and only borderline against full hybrid ($+0.058$, 95% CI $[0.000, 0.147]$, touching zero at the lower bound). Notably, the point estimate against full hybrid is now positive, not negative as in the raw descriptive comparison or the previous matched-subset measurement – consistent with, not contradicting, our overall reading that the raw descriptive gap in Table 12 likely does not reflect a real, consistent effect, though we no longer describe every comparison as uniformly non-significant, since two of the four now clear that bar in the expected direction. Sixth, this revision now runs on GPU (an NVIDIA RTX 3050 Ti, CUDA 12.1) rather than CPU-only, which resolves the previous hardware limitation for embedding fine-tuning and reranking, though BERTScore and generation-side wall-clock timings in Section 4.11 predate this change and were not re-measured. Seventh, the independent NLI-based faithfulness cross-check (Section 4.8.1) itself has weak, and on re-measurement chance-level, row-level agreement with the LLM judge (Pearson $r = -0.063$ as of the 2026-08-01 re-measurement, down from $r = 0.089$); we no longer treat this cross-check as corroborating evidence in either direction (Section 4.8.1 explains why), and note that building a genuinely reliable independent faithfulness signal remains open – to that end we prepared, but have not yet completed, a 50-item disagreement-prioritized human annotation sample (the query_ids where the LLM judge and the NLI check disagree most, plus a smaller agreement-anchor subsample), with blank rating columns for a human rater; this closes the machine-side half of a future human validation pass but the actual annotation, requiring a person, remains outside this revision’s scope. Eighth, the FacultyAvailability coverage gap identified in an earlier pass has been closed (Section 4.6), and we generalized the underlying lesson into a standing table-coverage audit (Section 3.10) rather than leaving it as two one-off fixes; the audit currently finds no remaining zero-coverage table, though it can only ever prove a table has *some* exercising query, not that its routing is correct in every case. Ninth, the LLM-based entity-normalization fallback (Section 3.7) implemented to address the entity-heavy abstention gate’s brittleness is not yet validated by its own ablation – we confirmed the specific spacing/dash gap that motivated it with a direct, deterministic regex fix instead (verified live), and left the broader LLM-based fallback available but disabled by default, the same honest posture as query translation before it was properly tested. Tenth, the expanded cross-lingual stress test (27 queries, three LLM-generated Banglish phrasings each of the same 9 underlying facts) shows a more cautious picture than the original 9-query set, and more cautious still on re-measurement: the direction still nominally favors query translation (59.3% vs. 63.0% top-5 accuracy as of 2026-08-01, both figures higher than previously measured because the underlying retriever itself improved) but is not statistically significant on the discordant pairs (exact binomial $p = 1.0$, up from $p = 0.219$), and manual inspection found visible naturalness problems in a minority of the higher-temperature-generated variants, consistent with Doğruöz et al. (2023)’s caution against treating LLM-generated code-switched text as equivalent to organically-collected data. The retriever-wide improvements made elsewhere this revision (BM25 retuning, the exact-match fix, the embedding retrain) have the side

effect of narrowing exactly the gap this stress test targets, on both the original and expanded sets – a real, disclosed interaction between unrelated fixes and this specific evaluation, not a sign the fixes were unhelpful.

5.2 Future Work

The most direct remaining priority is a human hallucination annotation study on the current pipeline, with a second rater, to complement the automated faithfulness score. This is, if anything, more urgent than before: an independent, architecturally-separate NLI-based faithfulness cross-check (Section 4.8.1), re-measured this revision, no longer cleanly corroborates the claim-decomposition-artifact hypothesis the way the previous measurement did – row-level agreement with the LLM judge is now indistinguishable from chance, so neither confirms nor rules out that the LLM-judge gap is a genuine grounding regression. A human rater remains the direct way to settle this, and the automated cross-check’s own instability under re-measurement is itself an argument for why an automated proxy alone should not be trusted to close this question. A second priority is extending the cross-lingual stress test’s small sample (Section 4.9.4) with more EnglishQA-only content as the corpus grows, since its current 9-query size limits it to a suggestive rather than a confirmatory result, even though its direction (query translation roughly doubling retrieval accuracy in the scenario it targets) is now demonstrated rather than merely hypothesized. A third is re-testing the sufficient-context override’s net effect on precision/recall now that the underlying confidence gate is substantially more sensitive (Table 3) than when the override was first validated.

This paper extended a hybrid BM25/dense RAG chatbot for academic advising into an adaptive pipeline combining domain-fine-tuned bilingual embeddings, per-query-type fusion routing, prerequisite-graph augmentation, and confidence-gated abstention with a sufficient-context override, and evaluated it on 200 English and 100 Banglish test queries with paired significance testing throughout, plus standard IR metrics (Recall@k, MRR, nDCG@10) and an independent NLI-based faithfulness cross-check. Under the current retriever, the adaptive pipeline is in full statistical parity with BM25-only on all four automated metrics – tied on BLEU/ROUGE-L/BERTScore, and only marginally, inconclusively behind on METEOR (significant under a paired t -test alone, not a matched Wilcoxon test) – and remains statistically tied with full hybrid and vector-only on all four as well; a deconfounding diagnostic found the router’s original fusion-method choice (RRF for entity-heavy queries) significantly negative in isolation once the exact-match mechanism’s ranking ceiling was removed, so the observed parity with the fixed baselines was borrowed entirely from that separate, independently-validated mechanism rather than being evidence that choice was beneficial – we replaced it with linear fusion at the same weighting as a direct, verified consequence, preserving the tie exactly while narrowing a residual ranking-quality gap. The system’s real, defensible value is the abstention and multi-hop graph reasoning capabilities the baselines cannot perform at all, together with the entity-collision handling that drives the exact-match mechanism’s own positive result, not a generation-quality edge over the single-method baselines. Embedding fine-tuning, an isolated unambiguous-match guarantee (extended in this revision to cover faculty lookups, not just course codes), and prerequisite-graph augmentation (confirmed on all four metrics, both significance tests, once measured on an expanded, still entirely real, $n = 31$ sample) each produced a measured positive effect; a cross-encoder reranker – both a generic version and, after we caught and corrected a training/deployment mismatch, a properly fine-tuned one – produced a measured negative effect in both cases; cross-lingual query translation showed a null effect on the main Banglish test set but a positive, now twice-replicated effect (61.5% to 84.6% top-5 accuracy at $n = 13$, narrower than first measured after this revision’s other retriever fixes partly closed the same gap) on a dedicated stress test built for the scenario it targets, though this effect no longer reaches significance on a larger, noisier expanded version of the same stress test; and BM25-level code-mixed normalization remains a null result on the data tested. We describe the result as a validated system-level integration for an under-studied bilingual advising setting, disclose every limitation and every methodological pitfall (including our own) found along the way, and identify human hallucination annotation as the most direct remaining step toward a stronger generalization claim.

ETHICAL STATEMENT

To protect student privacy, personal identifiers (names and student IDs) were removed from the dataset during ingestion, and community-sourced data was anonymized prior to analysis.

ACKNOWLEDGEMENTS

- **Conflicts of Interest:** The authors have no conflicts of interest or competing interests to declare in the context of this paper.
- **Financial and other Support:** No funding was received for the work reported here.
- **Author Contributions:** The authors contributed equally.

- **Data Sharing:** A dataset was generated and analyzed during the current study, available at <url>.
- **AI Tools:** The paper used AI tools for text correction purposes only.

References

- Adamopoulou, E. and Moussiades, L. (2021). Chatbots: History, technology, and applications. *Machine Learning with Applications*, 2:100006.
- Banerjee, S. and Lavie, A. (2005). METEOR: An automatic metric for MT evaluation with improved correlation with human judgments. In *Proceedings of ACL 2005*.
- Brandtzaeg, P. B. and Følstad, A. (2017). Why people use chatbots. In *Internet Science (INSCI 2017)*, LNCS vol. 10673, pages 377–392. Springer, Cham.
- Cormack, G. V., Clarke, C. L. A., and Buettcher, S. (2009). Reciprocal rank fusion outperforms Condorcet and individual rank learning methods. In *Proceedings of the 32nd International ACM SIGIR Conference*, pages 758–759.
- Athanasiou, D., Lymperaiou, M., Filandrianos, G., Voulodimos, A., and Stamou, G. (2026). AILS-NTUA at SemEval-2026 Task 8: Evaluating multi-turn RAG conversations. In *Proceedings of SemEval-2026*. arXiv:2603.10524.
- Huang, Y. and Huang, J. X. (2026). A survey on retrieval-augmented text generation for large language models. *ACM Computing Surveys*, 58(12), Article 300. arXiv:2404.10981.
- Dror, R., Baumer, G., Shlomov, S., and Reichart, R. (2018). The hitchhiker’s guide to testing statistical significance in natural language processing. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics*, pages 1383–1392.
- Doğruöz, A. S., Sitaram, S., and Yong, Z. X. (2023). Representativeness as a forgotten lesson for multilingual and code-switched data collection and preparation. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 5751–5767. arXiv:2310.20470.
- El-Ashmawi, W. H. et al. (2023). An interactive chatbot for college enquiry. *Journal of Computing and Communication*, 2(1):20–28.
- Es, S., James, J., Espinosa-Anke, L., and Schockaert, S. (2023). RAGAS: Automated evaluation of retrieval augmented generation. *arXiv preprint arXiv:2309.15217*.
- Formica, A., Mele, I., and Taglino, F. (2023). A template-based approach for question answering over knowledge bases. *Knowledge and Information Systems*, 65:453–479.
- Henderson, M., Al-Rfou, R., Strope, B., Sung, Y.-H., Lukacs, L., Guo, R., Kumar, S., Miklos, B., and Kurzweil, R. (2017). Efficient natural language response suggestion for smart reply. *arXiv preprint arXiv:1705.00652*.
- Goswami, S. and Kurra, S. (2025). HALT-RAG: A task-adaptable framework for hallucination detection with calibrated NLI ensembles and abstention. *arXiv preprint arXiv:2509.07475*.
- Laban, P., Schnabel, T., Bennett, P. N., and Hearst, M. A. (2022). SummaC: Re-visiting NLI-based models for inconsistency detection in summarization. *Transactions of the Association for Computational Linguistics*, 10:163–177.
- Hwang, T., Cho, S., Jeong, S., Song, H., Han, S., and Park, J. C. (2025). EXIT: Context-aware extractive compression for enhancing retrieval-augmented generation. In *Findings of the Association for Computational Linguistics: ACL 2025*. arXiv:2412.12559.
- Jin, B., Yoon, J., Han, J., and Arık, S. Ö. (2024). Long-context LLMs meet RAG: Overcoming challenges for long inputs in RAG. *arXiv preprint arXiv:2410.05983*.
- Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., and Liang, P. (2024). Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173.
- Magomere, J., La Malfa, E., Tonneau, M., Kazemi, A., and Hale, S. (2025). When claims evolve: Evaluating and enhancing the robustness of embedding models against misinformation edits. In *Findings of the Association for Computational Linguistics: ACL 2025*. arXiv:2503.03417.
- Meghwani, H., Agarwal, A., Pattanayak, P., Patel, H. L., and Panda, S. (2025). Hard negative mining for domain-specific retrieval in enterprise systems. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics: Industry Track*. arXiv:2505.18366.
- Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., et al. (2023). A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. *ACM Transactions on Information Systems*.
- Hwang, G.-J. and Chang, C.-Y. (2021). A review of opportunities and challenges of chatbots in education. *Interactive Learning Environments*, 29(1):1–14.

- Joren, H., Zhang, J., Ferng, C.-S., Juan, D.-C., Taly, A., and Rashtchian, C. (2024). Sufficient context: A new lens on retrieval augmented generation systems. *arXiv preprint arXiv:2411.06037*.
- Khanuja, S., Dandapat, S., Srinivasan, A., Sitaram, S., and Choudhury, M. (2020). GLUECoS: An evaluation benchmark for code-switched NLP. In *Proceedings of ACL 2020*.
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., et al. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33:9459–9474.
- Lin, C.-Y. (2004). ROUGE: A package for automatic evaluation of summaries. *Proceedings of ACL 2004*.
- Nguyen, H., Lopez, J., Homer, B., Ali, A., and Ahn, J. (2023). Reminders, reflections, and relationships: Insights from the design of a chatbot for college advising. *Information and Learning Sciences*, 124(3/4):128–146.
- Papineni, K., Roukos, S., Ward, T., and Zhu, W.-J. (2002). BLEU: A method for automatic evaluation of machine translation. In *Proceedings of ACL 2002*.
- Reimers, N. and Gurevych, I. (2019). Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In *Proceedings of EMNLP-IJCNLP 2019*.
- Rahman, M., Abedin, M., Abir, M. Z., Ansari, M. F. I., Reza, A., Sadeque, F. Y., and Farhan, N. (2025). Transforming mentorship: An AI powered chatbot approach to university guidance. *arXiv preprint arXiv:2511.04172*.
- Zeng, Q., Lu, Y., Zhou, Z., Qi, H., Yu, P., Zhao, F., Yanaka, H., Xuan, W., and Yokoya, N. (2026). Code-switching information retrieval: Benchmarks, analysis, and the limits of current retrievers. In *Findings of the Association for Computational Linguistics: ACL 2026*. arXiv:2604.17632.
- Kruk, F., Herath, S., and Choudhury, P. (2025). BanglAssist: A Bengali-English generative AI chatbot for code-switching and dialect-handling in customer service. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25)*. arXiv:2503.22283.
- Xiang, Z., Wu, C., Zhang, Q., Chen, S., Hong, Z., Huang, X., and Su, J. (2025). When to use graphs in RAG: A comprehensive analysis for graph retrieval-augmented generation. *arXiv preprint arXiv:2506.05690*.
- Jacob, M., Lindgren, E., Zaharia, M., Carbin, M., Khattab, O., and Drozdov, A. (2024). Drowning in documents: Consequences of scaling reranker inference. *ReNeuIR 2025 Workshop at SIGIR 2025*. arXiv:2411.11767.
- Chen, H., Long, Q., Xiao, M., Luo, X., Ju, W., Wang, C., Wang, X., Zhou, Y., and Zhu, H. (2025). SciRerankBench: Benchmarking rerankers towards scientific retrieval-augmented generated LLMs. *arXiv preprint arXiv:2508.08742*.
- Panickssery, A., Bowman, S. R., and Feng, S. (2024). LLM evaluators recognize and favor their own generations. *arXiv preprint arXiv:2404.13076*.
- Robertson, S. and Zaragoza, H. (2009). The probabilistic relevance framework: BM25 and beyond. *Foundations and Trends in Information Retrieval*.
- Sawarkar, K., Mangal, A., and Solanki, S. R. (2024). Blended RAG: Improving RAG (retriever-augmented generation) accuracy with semantic search and hybrid query-based retrievers. *arXiv preprint arXiv:2404.07220*.
- Shawar, B. A. and Atwell, E. (2007). Chatbots: Are they really useful? *LDV Forum*, 22(1):29–49.
- Swacha, J. and Gracel, M. (2025). Retrieval-augmented generation (RAG) chatbots for education: A survey of applications. *Applied Sciences*, 15(8):4234.
- Varshney, N., Yao, W., Zhang, H., Chen, J., and Yu, D. (2023). A stitch in time saves nine: Detecting and mitigating hallucinations of LLMs by validating low-confidence generation. *arXiv preprint arXiv:2307.03987*.
- Zhang, T., Kishore, V., Wu, F., Weinberger, K. Q., and Artzi, Y. (2019). BERTScore: Evaluating text generation with BERT. In *Proceedings of ICLR 2019*.